

Applied Machine Learning

Instructor:

Jeffrey W. Miller
Assistant Professor of Biostatistics
Harvard T.H. Chan School of Public Health
Building 1 Room 419, 655 Huntington Ave, Boston MA 02115
jwmiller@hsph.harvard.edu

Office hours: TBD

TAs: TBD

Course info: 2018 Spring 1-2, 5 Credits, Thirty (30) 1.75-hour Sessions

Time and location: TBD

Purpose of the Course

Statistical machine learning is a collection of flexible tools and techniques for using data to construct prediction algorithms for decision problems and exploratory analysis. This 5-credit course will take place in the second semester of the Health Data Science Master's program, and will introduce students to the most important elements of machine learning.

The central theme of the course will be to ground the material in practical real-world data examples, in order to motivate the concepts and illustrate why and how the methods work. Some mathematical foundations will be covered, but the primary emphasis of the course will be on learning how to implement and use the methods, while gaining an intuitive understanding of them. Programming (in R) and case studies will be used throughout the course to provide hands-on training.

Course Structure

The course will consist of 10 modules, building up from foundations, through linear, non-linear, and nonparametric supervised methods, as well as covering unsupervised learning and Bayesian methods. Along the way, the students will learn how to evaluate model performance/accuracy, quantify uncertainty, and combine methods via ensembles. The students will gain hands-on experience implementing and applying the methods in lab exercises and homework programming assignments, while learning the basic mathematical foundations in homework problem sets. There will be 29 lectures (including some in-class labs), and student performance will be assessed via 10 homework assignments, a midterm exam, and a final exam.

In addition to homework assignments and exams, class attendance and thoughtful participation are important and will be reflected in part in the final grade. Informed student participation in classroom discussions is required of all students. Students are expected to behave professionally at all times, with courtesy towards other students, the TAs, and the instructor.

Course Materials

Electronic copies of course notes/slides, homework assignments, and data sets will be posted on the course website. The textbook for the course is:

(JWHT) Gareth James, Daniella Witten, Trevor Hastie, Robert Tibshirani. “An Introduction to Statistical Learning”, Springer Texts in Statistics. Electronic copy available for free at: <http://www-bcf.usc.edu/~gareth/ISL/>

A more in-depth treatment of the same material is provided in:

Trevor Hastie, Robert Tibshirani, Jerome Friedman. “Elements of Statistical Learning”, Springer Texts in Statistics. Electronic copy available for free at: <http://statweb.stanford.edu/~tibs/ElemStatLearn/>

Prerequisites

This course is intended for students in all departments from the School. One of the following courses must be taken previously: BST 260 or BST 210 or BST 232, or an equivalent course.

Course Outline

Module 1: Getting started

Lecture 1: Course introduction, Probability basics

Reading: Course notes

Lecture 2: Overview, What is Statistical Learning?

Reading: JWHT 1, 2.1

Lecture 3: Assessing model accuracy, R programming basics

Reading: JWHT 2.2-2.3

Module 2: Linear regression

Lecture 4: Simple linear regression

Reading: JWHT 3.1

Lecture 5: Multiple linear regression

Reading: JWHT 3.2-3.3

Lecture 6: Illustrations, Lab on linear regression

Reading: JWHT 3.4-3.6

Module 3: Classification

Lecture 7: Classification overview, Logistic regression

Reading: JWHT 4.1-4.3

Lecture 8: Lab on classification

Reading: JWHT 4.6

Module 4: Resampling methods

Lecture 9: Cross-validation and bootstrap

Reading: JWHT 5.1-5.2

Lecture 10: Lab on cross-validation and bootstrap

Reading: JWHT 5.3

Module 5: Linear model selection and regularization

Lecture 11: Subset selection, Penalty-based methods

Reading: JWHT 6.1-6.2

Lecture 12: Lab on subset selection, ridge regression, and Lasso

Reading: JWHT 6.5-6.6

Lecture 13: Dimensional reduction, High-dimensional issues

Reading: JWHT 6.3-6.4

Lecture 14: Lab on dimensional reduction

Reading: JWHT 6.7

Midterm exam

Module 6: Non-linear regression

Lecture 15: Polynomial regression, step functions, and basis functions

Reading: JWHT 7.1-7.3

Lecture 16: Local regression, Generalized additive models (GAMs)

Reading: JWHT 7.6-7.7

Lecture 17: Lab on non-linear modeling

Reading: JWHT 7.8

Module 7: Tree-based methods and ensembles

Lecture 18: Classification and regression trees (CART), bagging, and random forests

Reading: JWHT 8.1-8.2

Lecture 19: Lab on CART and random forests

Reading: JWHT 8.3

Lecture 20: Boosting, gradient boosting, and other ensemble methods

Reading: JWHT 8.2, course notes

Lecture 21: Lab on ensemble methods

Reading: Course notes

Module 8: Support vector machines

Lecture 22: Maximum margin classifier, Support vector classifier

Reading: JWHT 9.1-9.2

Lecture 23: Kernel trick, Support vector machine, Multi-class problems

Reading: JWHT 9.3-9.4

Lecture 24: Lab on support vector machines

Reading: JWHT 9.6

Module 9: Unsupervised learning

Lecture 25: PCA review, K-means and hierarchical clustering

Reading: JWHT 10.1-10.3

Lecture 26: Lab on clustering

Reading: JWHT 10.5-10.6

Module 10: Bayesian methods

Lecture 27: Fundamentals of Bayesian statistics

Reading: Course notes

Lecture 28: Gibbs sampling

Reading: Course notes

Lecture 29: Lab on Gibbs sampling

Reading: Course notes

Final exam

Grading, Progress, and Performance

Performance in the course will be evaluated based on the following:

50% Homework

Homework assignments will consist of problem sets and computer programming assignments. The purpose of the homework assignments is to enable the students to more fully and deeply understand the concepts of the course, to gain practical experience implementing and using the methods introduced in the course, and to receive feedback on their performance and their understanding of the material.

There will 10 homework assignments, one for each module. Due dates will be posted and homework submission will be mandatory. Graded homework will be returned one week after submission.

Late submission policy: Homework submissions will be timestamped, and late submissions will be penalized as follows: starting from the due time until 24 hours after the due time, a multiplicative penalty starting at 1 and decreasing linearly to 0 will be applied. So, for example, an assignment submitted 6 hours late will incur a penalty of 0.75 (75% credit), and an assignment submitted 24 hours late will receive a penalty of 0.0 (no credit).

There will be no make-ups, extensions, or extra credit.

20% Midterm exam

25% Final exam

There will be one midterm exam and one final exam. The exams will consist of problems similar to those encountered by students in the homework assignments and the in-class exercises. The purpose of the exams is to evaluate the students understanding of the material, and provide feedback on their performance.

Exams will be closed book, closed notes. Exam dates will be announced far in advance, and there will be no make-up exams. Exams will be graded and returned within one week.

5% Class participation

Class attendance and thoughtful participation are important and will be reflected in part in the final grade. Informed student participation in classroom discussions is required of all students.

Harvard Chan Policies and Expectations

Inclusivity Statement

Diversity and inclusiveness are fundamental to public health education and practice. It is a requirement that you have an open mind and respect differences of all kinds. I share responsibility with you for creating a learning climate that is hospitable to all perspectives and cultures; please contact me if you have any concerns or suggestions.

Academic Integrity

Harvard University provides students with clear guidelines regarding academic standards and behavior. All work submitted to meet course requirements is expected to be a student's own work. In the preparation of work submitted to meet course requirements, students should always take great care to distinguish their own ideas and knowledge from information derived from sources. Please refer to [policy](#) in the student handbook for details on attributing credit and for doing independent work when required by the instructor.

Accommodations for Students with Disabilities

Harvard University provides academic accommodations to students with disabilities. Any requests for academic accommodations should ideally be made before the first week of the semester, except for unusual circumstances, so arrangements can be made. Students must register with the Local Disability Coordinator in the Office for Student Affairs to verify their eligibility for appropriate accommodations. Contact the OSA studentaffairs@hsph.harvard.edu in all cases, including temporary disabilities.

Course Evaluations

Constructive feedback from students is a valuable resource for improving teaching. The feedback should be specific, focused and respectful. It should also address aspects of the course and teaching that are positive as well as those which need improvement.

Completion of the evaluation is a requirement for each course. Your grade will not be available until you submit the evaluation. In addition, registration for future terms will be blocked until you have completed evaluations for courses in prior terms.