Asymptotic normality, concentration, and coverage of generalized posteriors

Jeffrey W. Miller

Harvard University, Department of Biostatistics

May 11, 2020

Abstract

Generalized likelihoods are commonly used to obtain consistent estimators with attractive computational and robustness properties. Formally, any generalized likelihood can be used to define a generalized posterior distribution, but an arbitrarily defined "posterior" cannot be expected to appropriately quantify uncertainty in any meaningful sense. In this article, we provide sufficient conditions under which generalized posteriors exhibit concentration, asymptotic normality (Bernstein–von Mises), an asymptotically correct Laplace approximation, and asymptotically correct frequentist coverage. We apply our results in detail to generalized posteriors for a wide array of generalized likelihoods, including pseudolikelihoods in general, the Ising model pseudolikelihood, the Gaussian Markov random field pseudolikelihood, the fully observed Boltzmann machine pseudolikelihood, the Cox proportional hazards partial likelihood, and a median-based likelihood for robust inference of location. Further, we show how our results can be used to easily establish the asymptotics of standard posteriors for exponential families and generalized linear models. We make no assumption of model correctness so that our results apply with or without misspecification.

Keywords: Bayesian theory, Bernstein–von Mises, composite likelihood, consistency, misspecification, pseudolikelihood, robustness.

1   Introduction

Many statistical estimation methods are based on maximizing a generalized likelihood function such as a pseudolikelihood, partial likelihood, or composite likelihood. Generalized likelihood functions are often advantageous in terms of computation or robustness while still having consistency guarantees, even though they do not necessarily correspond to the standard likelihood of a probabilistic model.

Formally, any generalized likelihood can be used to construct a generalized posterior proportional to the generalized likelihood times a prior. Generalized posteriors have been proposed based on a variety of generalized likelihoods, including composite likelihoods (Smith and Stephenson, 2009; Pauli et al., 2011; Ribatet et al., 2012; Friel, 2012), restricted likelihoods (Pettitt, 1983; Doksum and Lo, 1990; Hoff, 2007; Lewis et al., 2014), partial likelihoods

1

(Raftery et al., 1996; Sinha et al., 2003; Kim and Kim, 2009; Ventura and Racugno, 2016), substitution likelihoods (Lavine, 1995; Dunson and Taylor, 2005), modular likelihoods (Liu et al., 2009; Jacob et al., 2017), quasi-likelihoods (Ventura et al., 2010), generalized method of moments likelihoods (Yin, 2009), loss-based likelihoods (Jiang and Tanner, 2008; Zhang, 2006; Bissiri et al., 2016), and more. Although various theoretical guarantees have been provided for them, generalized posteriors have not yet been widely adopted, perhaps due to questions regarding their theoretical validity.

In this article, we provide new theoretical results on the asymptotic validity of generalized posteriors. We provide a range of sufficient conditions for concentration (Section 2), Bernstein–von Mises asymptotic normality and the Laplace approximation (Section 3), and asymptotic frequentist coverage of credible sets (Section 4) for generalized posteriors. For generalized posteriors derived from composite likelihoods—a large class covering essentially all the examples in this article—we informally discuss what can be expected in terms of consistency and coverage (Section 5). We show how our results can easily be applied to many standard posteriors, including i.i.d. exponential family models and (non-i.i.d.) generalized linear models for regression (Section 6). We then apply our results to generalized posteriors for an array of generalized likelihoods, including pseudolikelihoods in general, the Ising model pseudolikelihood, the Gaussian Markov random field pseudolikelihood, the fully observed Boltzmann machine pseudolikelihood, the Cox proportional hazards partial likelihood, and a median-based likelihood for robust inference of location (Section 7).

1.1 Novelty and comparison with previous work

In some sense, new Bernstein–von Mises (BvM) theorems are never surprising since they only verify what we already expect to happen if things are sufficiently nice. Thus, the utility of a BvM result is directly related to the ease and generality with which it can be applied. The main novelty of this article is that we provide results that are not only general, but are also relatively easy to apply in practice.

More specifically, the results in this article are novel in the following respects: (a) we provide rigorous results on generalized posteriors for non-i.i.d. data without any assumption of model correctness (in fact, in our main results, we do not even require there to be a probability model – true or assumed), (b) we provide sufficient conditions that are relatively easy to verify when they hold, and (c) we apply our results to a number of non-trivial examples, providing precise and concrete sufficient conditions for each example.

Standard BvM theorems are only applicable to standard posteriors under correctly specified i.i.d. probabilistic models (Van der Vaart, 2000; Ghosh and Ramamoorthi, 2003). Kleijn and Van der Vaart (2012) generalize by establishing a Bernstein–von Mises theorem under misspecification, but their result still only applies to standard posteriors, and they focus almost exclusively on the i.i.d. case. In contrast, our main results in Sections 2 and 3 do not involve a probability model at all and are applicable to arbitrary distributions of the form $\pi_n(\theta) \propto \exp(-nf_n(\theta))\pi(\theta)$, where the sequence of functions f_n is required to satisfy certain conditions. By treating the problem in this generality, we are able to provide results for i.i.d. and non-i.i.d. cases with or without misspecification; see the examples in Sections 6 and 7. Additionally, BvM theorems often only show that the total variation distance converges to zero in probability; in contrast, we prove it converges to zero almost surely.

For semiparametric and nonparametric models, a number of BvM results have been established (Shen, 2002; Kim and Lee, 2004; Leahu, 2011; Castillo and Nickl, 2013; Bickel and Kleijn, 2012; Castillo and Rousseau, 2013). Here, we focus on the parametric case in which θ has fixed, finite dimension—however, if θ is a finite-dimensional functional of a semiparametric or nonparametric model, then in principle our main results could still be applied to the posterior of θ since our conditions are stated directly in terms of f_n rather than in terms of a probability model.

A very general BvM result is provided by Panov and Spokoiny (2015), who establish a finite-sample BvM theorem for non-i.i.d. semiparametric models under misspecification, allowing the dimension of the parameter to grow with the sample size. While their results are very general, their conditions are quite abstract and seem difficult to verify, particularly for non-experts. Bochkina and Green (2014) provide an interesting BvM result that handles cases in which the true parameter is on the boundary of the parameter space, and their result is also applicable when the assumed probability model is misspecified.

For generalized posteriors, much of the previous work on asymptotic normality tends to rely on unspecified regularity conditions or only establishes weak convergence, that is, convergence in distribution (Doksum and Lo, 1990; Lazar, 2003; Greco et al., 2008; Pauli et al., 2011; Ribatet et al., 2012; Ventura and Racugno, 2016). In contrast, we show convergence in total variation distance and we provide rigorous results with all assumptions explicitly stated. Further, the usual regularity conditions in previous work include an assumption of concentration (Bernardo and Smith, 2000); in contrast, we prove concentration.

In general, we make no assumption of model correctness. However, to ensure that a generalized posterior is doing something reasonable, it is desirable to have a guarantee of consistency—that is, concentration at the true parameter—if the assumed model is correct or at least partially correct. To this end, in Section 5 we show that for any composite likelihood derived from a correct model, the resulting generalized posterior concentrates at the true parameter under fairly general conditions. Since many generalized likelihoods can be viewed as composite likelihoods, this establishes consistency in a wide range of cases. On the other hand, it is well-known that—except in special circumstances—the asymptotic frequentist coverage of composite likelihood-based posteriors is typically incorrect unless an adjustment is made (Pauli et al., 2011; Ribatet et al., 2012); see Section 5 for more details.

For each main result in Sections 2 and 3, we provide a range of alternative sufficient conditions, from more abstract to more concrete. The more abstract versions are more generally applicable, whereas the more concrete versions have conditions that are easier to verify when applicable. For instance, Theorem 3.1 is an abstract BvM theorem involving a quadratic representation condition; meanwhile, Theorem 3.2 is a more concrete BvM theorem involving conditions on derivatives that are roughly analogous to the conditions in classical BvM theorems (Ghosh and Ramamoorthi, 2003, Theorem 1.4.2). We also provide versions of the theorems based on convexity of f_n (see Theorems 2.3(3) and 3.2(2)), which is usually easy to verify when it applies and simplifies the other required conditions; this is very roughly analogous to convexity-based results on asymptotic normality of estimators (Hjort and Pollard, 1993).

2 Posterior concentration

Theorem 2.2 is a general concentration result for generalized posteriors Π_n on a measurable space $(\Theta, \mathcal{A})$. The basic structure of the proof of Theorem 2.2 follows that of Schwartz's theorem (Schwartz, 1965; Ghosh and Ramamoorthi, 2003). Although Theorem 2.2 is useful for theoretical purposes, in practice, one typically needs to establish concentration on neighborhoods in a relevant topology on Θ. To this end, Theorem 2.3 provides a range of sufficient conditions for concentration on metric space neighborhoods of a point $\theta_0 \in \Theta$.

Condition 2.1. *Let $f_n : \Theta \to \mathbb{R}$ for $n \in \mathbb{N}$ be a sequence of functions on a probability space $(\Theta, \mathcal{A}, \Pi)$. For all n, assume $z_n < \infty$ where $z_n = \int_\Theta \exp(-nf_n(\theta))\Pi(d\theta)$, and define the probability measure*

$$\Pi_n(d\theta) = \exp(-nf_n(\theta))\Pi(d\theta)/z_n.$$

Throughout, all arbitrarily defined functions and sets are assumed to be measurable, and we denote $\mathbb{N} = \{1, 2, \ldots\}$. Here, $\exp(-nf_n(\theta))$ is interpreted as the "likelihood", possibly in some generalized sense, Π is the "prior", and Π_n is the "posterior".

We reemphasize that our main theorems in Sections 2 and 3 do not involve a probability model and do not even require that there be data. Instead, our results apply to arbitrary deterministic sequences of distributions Π_n satisfying certain conditions. Consequently, the mode of convergence in these theorems is not probabilistic in any sense. In the applications in Sections 6 and 7 that involve probability models, we show that the conditions apply with probability 1, and in this way we obtain almost sure convergence.

Theorem 2.2. *Assume Condition 2.1. If $\theta_0 \in \Theta$ and there exists $f : \Theta \to \mathbb{R}$ such that*

1. *$f_n(\theta) \to f(\theta)$ as $n \to \infty$ for all $\theta \in \Theta$,*

2. *$\Pi(A_\varepsilon) > 0$ for all $\varepsilon > 0$, where $A_\varepsilon = \{\theta \in \Theta : f(\theta) < f(\theta_0) + \varepsilon\}$, and*

3. *$\liminf_n \inf_{\theta \in A_\varepsilon^c} f_n(\theta) > f(\theta_0)$ for all $\varepsilon > 0$,*

then $\Pi_n(A_\varepsilon) \to 1$ as $n \to \infty$, for any $\varepsilon > 0$.

See Section S1 for the proof. When Θ is a metric space, the collection of functions (f_n) is said to be *equicontinuous* if for any $\varepsilon > 0$ there exists $\delta > 0$ such that for all $n \in \mathbb{N}$, $\theta, \theta' \in \Theta$, if $d(\theta, \theta') < \delta$ then $|f_n(\theta) - f_n(\theta')| < \varepsilon$. For a function $f : E \to \mathbb{R}$ where $E \subseteq \mathbb{R}^D$, we denote the gradient by $f'(\theta)$ (that is, $f'(\theta) = (\frac{\partial f}{\partial \theta_i}(\theta))_{i=1}^D \in \mathbb{R}^D$) and the Hessian by $f''(\theta)$ (that is, $f''(\theta) = (\frac{\partial^2 f}{\partial \theta_i \partial \theta_j}(\theta))_{i,j=1}^D \in \mathbb{R}^{D \times D}$) when these derivatives exist. We use the following definition of convexity to allow the possibility that the domain $E \subseteq \mathbb{R}^D$ is non-convex: $f : E \to \mathbb{R}$ is *convex* if for all $\theta, \theta' \in E$ and all $t \in [0, 1]$ such that $t\theta + (1-t)\theta' \in E$, we have $f(t\theta + (1-t)\theta') \leq tf(\theta) + (1-t)f(\theta')$.

Theorem 2.3. *Assume Condition 2.1. Suppose (Θ, d) is a metric space and $\mathcal{A}$ is the resulting Borel sigma-algebra. Let $\theta_0 \in \Theta$ and denote $N_\varepsilon = \{\theta \in \Theta : d(\theta, \theta_0) < \varepsilon\}$. If $\Pi(N_\varepsilon) > 0$ for all $\varepsilon > 0$, $f_n \to f$ pointwise on Θ for some $f : \Theta \to \mathbb{R}$, and any one of the following three sets of conditions hold, then for any $\varepsilon > 0$, $\Pi_n(N_\varepsilon) \to 1$ as $n \to \infty$.*

1. f is continuous at θ_0 and $\liminf_n \inf_{\theta \in N_\varepsilon^c} f_n(\theta) > f(\theta_0)$ for all $\varepsilon > 0$.

2. (f_n) is equicontinuous on some compact set $K \subseteq \Theta$, θ_0 is an interior point of K, $f(\theta) > f(\theta_0)$ for all $\theta \in K \setminus \{\theta_0\}$, and $\liminf_n \inf_{\theta \in K^c} f_n(\theta) > f(\theta_0)$.

3. f_n is convex for each n, $\Theta \subseteq \mathbb{R}^D$ with the Euclidean metric, θ_0 is an interior point of Θ, and either

   (a) $f(\theta) > f(\theta_0)$ for all $\theta \in \Theta \setminus \{\theta_0\}$, or

   (b) f' exists in a neighborhood of θ_0, $f'(\theta_0) = 0$, and $f''(\theta_0)$ exists and is positive definite.

Further, 2 ⇒1 and 3 ⇒1 under the assumptions of the theorem.

See Section S1 for the proof. Note that if Θ is compact, then case 2 with $K = \Theta$ simplifies to (f_n) being equicontinuous and $f(\theta) > f(\theta_0)$ for all $\theta \in \Theta \setminus \{\theta_0\}$. This can be used to prove consistency results based on classical Wald-type conditions such as in Ghosh and Ramamoorthi (2003) 1.3.4.

3 Asymptotic normality and Laplace approximation

Theorem 3.1 establishes general sufficient conditions under which a generalized posterior exhibits asymptotic normality and an asymptotically correct Laplace approximation, along with concentration at θ_0. As in Section 2, $\pi(\theta)$ can be interpreted as the prior density and $\pi_n(\theta) \propto \exp(-n f_n(\theta)) \pi(\theta)$ can be thought of as the "posterior" density. The points θ_n can be viewed as maximum generalized likelihood estimates. The proof of Theorem 3.1 is concise, but some of the conditions of the theorem are a bit abstract. Thus, we also provide Theorem 3.2 to give more concrete sufficient conditions which, when satisfied, are usually easier to verify. Theorem 3.2 is the main result used in the examples in the rest of the paper.

Unlike previous work on BvM, the results in this section only involve conditions on f_n and π, and do not involve any assumptions at all regarding the data; indeed, the results in this section and Section 2 do not even require that there be any data. We also highlight two supporting results that are employed in the proof of Theorem 3.2. Theorem 3.3 provides concrete sufficient conditions under which the quadratic representation (condition 1) in Theorem 3.1 holds. Theorem 3.4 is a purely analytic result on uniform convergence of f_n, f_n', and f_n'', which we believe is interesting in its own right.

Given $x_0 \in \mathbb{R}^D$ and $r > 0$, we write $B_r(x_0)$ to denote the open ball of radius r at x_0, that is, $B_r(x_0) = \{x \in \mathbb{R}^D : |x - x_0| < r\}$. We use $|\cdot|$ to denote the Euclidean norm. Given positive sequences (a_n) and (b_n), we write $a_n \sim b_n$ to denote that $a_n/b_n \to 1$ as $n \to \infty$. We write $\mathcal{N}(x \mid \mu, C)$ to denote the normal density with mean μ and covariance matrix C.

Theorem 3.1. *Let $\theta_0 \in \mathbb{R}^D$ and let $\pi : \mathbb{R}^D \to \mathbb{R}$ be a probability density with respect to Lebesgue measure such that π is continuous at θ_0 and $\pi(\theta_0) > 0$. Let $f_n : \mathbb{R}^D \to \mathbb{R}$ for $n \in \mathbb{N}$ and assume:*

1. f_n can be represented as

$$f_n(\theta) = f_n(\theta_n) + \tfrac{1}{2}(\theta - \theta_n)^{\mathsf{T}} H_n(\theta - \theta_n) + r_n(\theta - \theta_n) \qquad (1)$$

where $\theta_n \in \mathbb{R}^D$ such that $\theta_n \to \theta_0$, $H_n \in \mathbb{R}^{D \times D}$ symmetric such that $H_n \to H_0$ for some positive definite H_0, and $r_n : \mathbb{R}^D \to \mathbb{R}$ has the following property: there exist $\varepsilon_0, c_0 > 0$ such that for all n sufficiently large, for all $x \in B_{\varepsilon_0}(0)$, we have $|r_n(x)| \le c_0 |x|^3$; and

2. for any $\varepsilon > 0$, $\liminf_n \inf_{\theta \in B_\varepsilon(\theta_n)^c} \left(f_n(\theta) - f_n(\theta_n) \right) > 0$.

Then, defining $z_n = \int_{\mathbb{R}^D} \exp(-n f_n(\theta)) \pi(\theta) d\theta$ and $\pi_n(\theta) = \exp(-n f_n(\theta)) \pi(\theta) / z_n$, we have

$$\int_{B_\varepsilon(\theta_0)} \pi_n(\theta) d\theta \xrightarrow[n \to \infty]{} 1 \text{ for all } \varepsilon > 0, \qquad (2)$$

that is, π_n concentrates at θ_0,

$$z_n \sim \frac{\exp(-n f_n(\theta_n)) \pi(\theta_0)}{|\det H_0|^{1/2}} \left(\frac{2\pi}{n} \right)^{D/2} \qquad (3)$$

as $n \to \infty$ (Laplace approximation), and letting q_n be the density of $\sqrt{n}(\theta - \theta_n)$ when $\theta \sim \pi_n$,

$$\int_{\mathbb{R}^D} \left| q_n(x) - \mathcal{N}(x \mid 0, H_0^{-1}) \right| dx \xrightarrow[n \to \infty]{} 0, \qquad (4)$$

that is, q_n converges to $\mathcal{N}(0, H_0^{-1})$ in total variation.

See Section S2 for the proof. Throughout, we use the Euclidean–Frobenius norms on vectors $v \in \mathbb{R}^D$, matrices $M \in \mathbb{R}^{D \times D}$, and tensors $T \in \mathbb{R}^{D^3}$, that is, $|v| = (\sum_i v_i^2)^{1/2}$, $\|M\| = (\sum_{i,j} M_{ij}^2)^{1/2}$, and $\|T\| = (\sum_{i,j,k} T_{ijk}^2)^{1/2}$. Convergence and boundedness for vectors, matrices, and tensors is defined with respect to these norms. A collection of functions $h_n : E \to F$, where F is a normed space, is *uniformly bounded* if the set $\{\|h_n(x)\| : x \in E, \, n \in \mathbb{N}\}$ is bounded, and is *pointwise bounded* if $\{\|h_n(x)\| : n \in \mathbb{N}\}$ is bounded for each $x \in E$. Let $f'''(\theta)$ denote the tensor of third derivatives, that is, $f'''(\theta) = (\frac{\partial^3 f}{\partial \theta_i \partial \theta_j \partial \theta_k}(\theta))_{i,j,k=1}^D \in \mathbb{R}^{D^3}$.

Theorem 3.2. *Let $\Theta \subseteq \mathbb{R}^D$. Let $E \subseteq \Theta$ be open (in $\mathbb{R}^D$), convex, and bounded. Let $\theta_0 \in E$ and let $\pi : \Theta \to \mathbb{R}$ be a probability density with respect to Lebesgue measure such that π is continuous at θ_0 and $\pi(\theta_0) > 0$. Let $f_n : \Theta \to \mathbb{R}$ have continuous third derivatives on E. Suppose $f_n \to f$ pointwise for some $f : \Theta \to \mathbb{R}$, $f''(\theta_0)$ is positive definite, and (f_n''') is uniformly bounded on E. If either of the following two conditions is satisfied:*

1. *$f(\theta) > f(\theta_0)$ for all $\theta \in K \setminus \{\theta_0\}$ and $\liminf_n \inf_{\theta \in \Theta \setminus K} f_n(\theta) > f(\theta_0)$ for some compact $K \subseteq E$ with θ_0 in the interior of K, or*

2. *each f_n is convex and $f'(\theta_0) = 0$,*

then there is a sequence $\theta_n \to \theta_0$ such that $f_n'(\theta_n) = 0$ for all n sufficiently large, $f_n(\theta_n) \to f(\theta_0)$, Equation 2 (concentration at θ_0) holds, Equation 3 (Laplace approximation) holds, and Equation 4 (asymptotic normality) holds, where $H_0 = f''(\theta_0)$. Further, condition 2 implies condition 1 under the assumptions of the theorem.

See Section S2 for the proof. The following is used in the proof of Theorem 3.2.

Theorem 3.3. *Let $E \subseteq \mathbb{R}^D$ be open and convex, and let $\theta_0 \in E$. Let $f_n : E \to \mathbb{R}$ have continuous third derivatives, and assume:*

1. *there exist $\theta_n \in E$ such that $\theta_n \to \theta_0$ and $f_n'(\theta_n) = 0$ for all n sufficiently large,*

2. *$f_n''(\theta_0) \to H_0$ as $n \to \infty$ for some positive definite H_0, and*

3. *(f_n''') is uniformly bounded.*

Then, letting $H_n = f_n''(\theta_n)$, condition 1 of Theorem 3.1 is satisfied for all n sufficiently large.

See Section S2 for the proof. The main tool used in the proof of Theorem 3.2 is the following result, which provides somewhat more than we require. A collection of functions $h_n : E \to F$, where E and F are subsets of normed spaces, is *equi-Lipschitz* if there exists $c > 0$ such that for all $n \in \mathbb{N}$, $x, y \in E$, we have $\|h_n(x) - h_n(y)\| \leq c\|x - y\|$.

Theorem 3.4 (Regular convergence). *Let $E \subseteq \mathbb{R}^D$ be open, convex, and bounded. For $n \in \mathbb{N}$, let $f_n : E \to \mathbb{R}$ have continuous third derivatives, and suppose (f_n''') is uniformly bounded. If (f_n) is pointwise bounded, then (f_n), (f_n'), and (f_n'') are all equi-Lipschitz and uniformly bounded. If $f_n \to f$ pointwise for some $f : E \to \mathbb{R}$, then f' and f'' exist, $f_n \to f$ uniformly, $f_n' \to f'$ uniformly, and $f_n'' \to f''$ uniformly.*

Note that if $f_n \to f$ pointwise then (f_n) is pointwise bounded; thus, if $f_n \to f$ pointwise then we also get the equi-Lipschitz and uniform bounded result. See Section S3 for the proof.

4 Coverage

For a generalized posterior to provide useful quantification of uncertainty, it is important that it be well-calibrated in terms of frequentist coverage. Ideally, we would like Π_n to have correct frequentist coverage in the sense that posterior credible sets of probability α have frequentist coverage α. Obviously, an arbitrarily chosen generalized posterior cannot be expected to have correct coverage. Thus, in Theorem 4.1, we provide simple sufficient conditions under which a generalized posterior has correct frequentist coverage, asymptotically.

To interpret Theorem 4.1, we think of θ_n as a maximum generalized likelihood estimate, θ_0 as the "true" parameter we want to cover, Π_n as the generalized posterior distribution, S_n as a credible set of asymptotic probability α, Q_n as a centered and scaled version of Π_n, and R_n as a centered and scaled version of S_n. Roughly, the theorem says that if Q_n converges in total variation to the asymptotic distribution of $-\sqrt{n}(\theta_n - \theta_0)$, and R_n converges pointwise, then asymptotically, S_n contains the true parameter 100α percent of the time. In other words, if the conditions of the theorem hold, then asymptotically, Π_n has correct frequentist coverage in the sense that posterior credible sets of probability α have frequentist coverage α.

Typically, when things work out nicely, θ_n is $\sqrt{n}$-consistent and asymptotically normal and a BvM result holds for Π_n, in which case the result says that Π_n has correct coverage asymptotically if the covariance matrices of these two normal distributions are equal. In other words, if $\sqrt{n}(\theta_n - \theta_0) \xrightarrow{\mathrm{D}} \mathcal{N}(0, C_1)$ and $Q_n \to \mathcal{N}(0, C_2)$ in total variation distance, then

Π_n has correct asymptotic frequentist coverage if $C_1 = C_2$ and the other conditions hold. In this case, the only other condition is that R_n converges to a set R with finite nonzero Lebesgue measure, because it is guaranteed that $Q(\partial R) = 0$. (Note that if $X \sim \mathcal{N}(0, C_1)$ then $-X \sim \mathcal{N}(0, C_1)$ also.) This result is precisely what one would expect; thus, the purpose of the theorem is to make this rigorous under easy-to-verify conditions.

In Section 5.2, we observe that composite likelihood-based generalized posteriors do not generally have correct frequentist coverage, except in special circumstances. We expect that many generalized posteriors (like misspecified models) do not have correct frequentist coverage, even asymptotically, since this is a fairly stringent requirement. For this reason, our results in Sections 6 and 7 do not include coverage guarantees. On the other hand, having correct coverage for each univariate component of the parameter, marginally, rather than jointly, is a much less stringent requirement, and for future work it might be interesting to investigate conditions under which this holds for the generalized posteriors we consider. Alternatively, one could aim for conservative coverage (that is, covering θ_0 with probability $\geq \alpha$), rather than exactly correct coverage, which would be more achievable as well.

We give $\mathbb{R}^D$ the Euclidean topology and the resulting Borel sigma-algebra, $\mathcal{B}$, and we use $m(\cdot)$ to denote Lebesgue measure on $\mathbb{R}^D$. We write ∂R to denote the boundary of a set $R \in \mathbb{R}^D$, that is, $\partial R = \bar{R} \setminus R^\circ$, where $\bar{R}$ is the closure and R° is the interior of R. Given $R, R_1, R_2, \ldots \subseteq \mathbb{R}^D$, we write $R_n \to R$ to denote that for all $x \in \mathbb{R}^D$, $\mathbb{1}(x \in R_n) \to \mathbb{1}(x \in R)$ as $n \to \infty$. Define $d(x, A) = \inf_{y \in A} \|x - y\|$ for $x \in \mathbb{R}^D$ and $A \subseteq \mathbb{R}^D$.

Theorem 4.1. *Let $\theta_1, \theta_2, \ldots \in \mathbb{R}^D$ be a sequence of random vectors, and let $\theta_0 \in \mathbb{R}^D$ be fixed. Let $\Pi_1, \Pi_2, \ldots$ be a sequence of random probability measures on $\mathbb{R}^D$, possibly dependent on $\theta_1, \theta_2, \ldots$. Let $S_1, S_2, \ldots \subseteq \mathbb{R}^D$ be a sequence of random convex measurable sets such that $\Pi_n(S_n) \xrightarrow{\text{a.s.}} \alpha$ for some fixed $\alpha \in (0, 1)$. For $A \in \mathcal{B}$, define $Q_n(A) = \int \mathbb{1}(\sqrt{n}(\theta - \theta_n) \in A)\Pi_n(d\theta)$ and define $R_n = \{\sqrt{n}(\theta - \theta_n) : \theta \in S_n\}$. Suppose there is a fixed probability measure Q and a fixed set $R \subseteq \mathbb{R}^D$ such that*

1. *$-\sqrt{n}(\theta_n - \theta_0) \xrightarrow{\text{D}} Q$ as $n \to \infty$ (where $\xrightarrow{\text{D}}$ denotes convergence in distribution),*

2. *$\sup_{A \in \mathcal{B}} |Q_n(A) - Q(A)| \xrightarrow{\text{a.s.}} 0$ as $n \to \infty$ (that is, $Q_n \xrightarrow{\text{a.s.}} Q$ in total variation),*

3. *$R_n \xrightarrow{\text{a.s.}} R$ as $n \to \infty$ (that is, a.s., for all $x \in \mathbb{R}^D$, $\mathbb{1}(x \in R_n) \to \mathbb{1}(x \in R)$), and*

4. *$Q(\partial R) = 0$ and $0 < m(R) < \infty$, where m denotes Lebesgue measure on $\mathbb{R}^D$.*

Then $\mathbb{P}(\theta_0 \in S_n) \to \alpha$ as $n \to \infty$.

See Section S4 for the proof. If Q has a density with respect to Lebesgue measure, then the condition that $Q(\partial R) = 0$ is automatically satisfied, since the assumptions imply that R is convex and thus $m(\partial R) = 0$. The following lemmas are used in the proof of Theorem 4.1, but may be useful in their own right.

Lemma 4.2. *Let $X_1, X_2, \ldots \in \mathbb{R}^D$ be random vectors such that $X_n \xrightarrow{\text{D}} X$ for some random vector X. Let $R_1, R_2, \ldots \subseteq \mathbb{R}^D$ be random convex measurable sets, possibly dependent on $X_1, X_2, \ldots$. Assume there exists some fixed $R \subseteq \mathbb{R}^D$ with $0 < m(R) < \infty$ and $\mathbb{P}(X \in \partial R) = 0$ such that $R_n \to R$ almost surely as $n \to \infty$. Then $\mathbb{P}(X_n \in R_n) \to \mathbb{P}(X \in R)$ as $n \to \infty$.*

8

See Section S4 for the proof. The probability $\mathbb{P}(X_n \in R_n)$ should be interpreted as $\int \mathbb{1}(X_n(\omega) \in R_n(\omega))P(d\omega)$, that is, X_n and R_n are jointly integrated over and $\mathbb{P}(X_n \in R_n)$ is a non-random quantity.

Lemma 4.3. *Let* $R_1, R_2, \ldots \subseteq \mathbb{R}^D$ *be convex sets. Assume* $R_n \to R$ *for some* $R \subseteq \mathbb{R}^D$ *with* $0 < m(R) < \infty$. *For any* $\varepsilon > 0$, *if* $A = \{x \in \mathbb{R}^D : d(x, R^c) > \varepsilon\}$ *and* $B = \{x \in \mathbb{R}^D : d(x, R) \leq \varepsilon\}$ *then for all* n *sufficiently large,* $A \subseteq R_n \subseteq B$.

See Section S4 for the proof.

5 Composite likelihood-based posteriors

Composite likelihoods (CLs) (Lindsay, 1988) represent a large class of generalized likelihoods that encompasses essentially all of the examples in Sections 6 and 7. The theory of maximum composite likelihood estimation is well-established (Lindsay, 1988; Molenbergs and Verbeke, 2005; Varin et al., 2011) and theoretical results for CL-based generalized posteriors have been provided (Pauli et al., 2011; Ribatet et al., 2012; Ventura and Racugno, 2016; Greco et al., 2008; Lazar, 2003). In this section, we informally discuss what can be expected of CL-based generalized posteriors, or CL-posteriors for short, based on our results in Sections 2-4. Roughly speaking, CL-posteriors derived from a correctly specified model can generally be expected to be consistent, but not necessarily correctly calibrated with respect to frequentist coverage. The consistency and coverage of CL-posteriors has been studied in previous work, subject to the caveats discussed in the introduction (Pauli et al., 2011; Ribatet et al., 2012). The purpose of this section is to illustrate how these previous results can be strengthened using our results in Sections 2-4.

Let y denote the full data set, which may take any form such as a sequence, a graph, a database, or any other data structure. Suppose $\{P_\theta : \theta \in \Theta\}$ is an assumed model for the distribution of y given θ, where $\Theta \subseteq \mathbb{R}^D$. For $j = 1, \ldots, k$, suppose $s_j(y)$ and $t_j(y)$ are functions of the data and, when $Y \sim P_\theta$, suppose the conditional distribution of $s_j(Y)$ given $t_j(Y)$ has density $p_\theta(s_j | t_j)$ with respect to a common dominating measure λ_j for all values of θ and t_j. Define the *composite likelihood* (Lindsay, 1988),

$$\mathcal{L}^{\mathrm{CL}}(\theta) = \prod_{j=1}^{k} p_\theta(s_j | t_j).$$

A few examples are given here and in Section 7; see Varin et al. (2011) for more examples.

Example 5.1 (i.i.d. likelihood)**.** If $y = (y_1, \ldots, y_n)$, $s_j(y) = y_j$, and $t_j(y) = 0$, then $\mathcal{L}^{\mathrm{CL}}(\theta) = \prod_{j=1}^{n} p_\theta(y_j)$ is simply the likelihood of an i.i.d. model.

Example 5.2 (pseudolikelihood)**.** If $y = (y_1, \ldots, y_n)$, $s_j(y) = y_j$, and $t_j(y) = y_{-j} := (y_1, \ldots, y_{j-1}, y_{j+1}, \ldots, y_n)$, then $\mathcal{L}^{\mathrm{CL}}(\theta) = \prod_{j=1}^{n} p_\theta(y_j | y_{-j})$ is a pseudolikelihood (Besag, 1975).

Example 5.3 (restricted likelihood)**.** If $k = 1$, $t_1(y) = 0$, and $s_1(y)$ is an insufficient statistic, then $\mathcal{L}^{\mathrm{CL}}(\theta)$ is a restricted likelihood (Lewis et al., 2014). For instance, if $s_1(y)$ consists of ranks or selected quantiles, then $\mathcal{L}^{\mathrm{CL}}(\theta)$ is a rank likelihood (Pettitt, 1983; Hoff, 2007) or a quantile-based likelihood (Doksum and Lo, 1990), respectively.

9

Due to the structure of composite likelihoods, one can make some general observations about CL-posteriors of the form $\pi_n(\theta) \propto \mathcal{L}^{\mathrm{CL}}(\theta)\pi(\theta)$. First, a reassuring property is that if the model is correctly specified, then CL-posteriors are consistent under fairly general conditions; we discuss this next.

5.1   Consistency of CL-posteriors under correct specification

Throughout this article, we make no assumption of model correctness in the main results (Sections 2-4) or the applications (Sections 6-7). However, for interpretability, it is important to have a guarantee of consistency if the assumed model is correct or at least partially correct. Here, we show that in many cases of interest, if the model is correctly specified—or at least, if the conditional densities $p_\theta(s_j|t_j)$ are correctly specified—then the CL-posterior concentrates at the true parameter. The analogue of this result for maximum CL estimators is well-known (Lindsay, 1988; Varin et al., 2011); also see Pauli et al. (2011) and Ribatet et al. (2012).

First, observe that if $Y \sim P_{\theta_0}$, $S_j = s_j(Y)$, and $T_j = t_j(Y)$, then for all $\theta \in \Theta$,

$$\mathrm{E}\big(\log p_{\theta_0}(S_j|T_j)\big) \geq \mathrm{E}\big(\log p_\theta(S_j|T_j)\big) \tag{5}$$

because the conditional relative entropy $\mathrm{E}\big(\log(p_{\theta_0}(S_j|T_j)/p_\theta(S_j|T_j))\big)$ is nonnegative; this is referred to as the information inequality by Lindsay (1988). Now, suppose that for each $n \in \{1, 2, \ldots\}$, we have a data set Y^n, model $\{P_\theta^n : \theta \in \Theta\}$ (where Θ does not depend on n), and functions s_j^n, t_j^n for $j = 1, \ldots, k_n$. Further, suppose the assumed model is correct, such that $Y^n \sim P_{\theta_0}^n$ where the true parameter θ_0 is shared across all n. Define

$$f_n(\theta) = -\frac{1}{n}\log \mathcal{L}_n^{\mathrm{CL}}(\theta) = -\frac{1}{n}\sum_{j=1}^{k_n}\log p_\theta^n(S_j^n|T_j^n)$$

and $\pi_n(\theta) \propto \exp(-nf_n(\theta))\pi(\theta) = \mathcal{L}_n^{\mathrm{CL}}(\theta)\pi(\theta)$ where $S_j^n = s_j^n(Y^n)$ and $T_j^n = t_j^n(Y^n)$. In many cases of interest (see Sections 6 and 7), we have that with probability 1, for all $\theta \in \Theta$, $\lim_{n\to\infty} f_n(\theta) = f(\theta)$ where $f(\theta) = \lim_{n\to\infty}\mathrm{E}\big(f_n(\theta)\big)$. Then, by Equation 5, $f(\theta_0) \leq f(\theta)$ for all $\theta \in \Theta$, in other words, θ_0 is a minimizer of f. Further, in many cases, f has a unique minimizer, and π_n concentrates at the unique minimizer; in particular, this holds if the conditions of Theorem 2.3 or Theorem 3.2 are met. Therefore, in such cases, the CL-posterior π_n concentrates at the true parameter, θ_0.

5.2   Coverage of CL-posteriors under correct specification

Although CL-posteriors have appealing consistency properties, they do not generally have correct asymptotic frequentist coverage, except in special circumstances (Pauli et al., 2011; Ribatet et al., 2012). Continuing in the notation of Section 5.1, suppose $Y^n \sim P_{\theta_0}^n$, let $\pi_n(\theta) \propto \exp(-nf_n(\theta))\pi(\theta) = \mathcal{L}_n^{\mathrm{CL}}(\theta)\pi(\theta)$ be the CL-posterior, and let $\theta_n = \arg\max_\theta \mathcal{L}_n^{\mathrm{CL}}(\theta) = \arg\min_\theta f_n(\theta)$ be the maximum composite likelihood estimator. If Theorem 3.2 applies with probability 1, then $Q_n \xrightarrow{\text{a.s.}} \mathcal{N}(0, H_0^{-1})$ in total variation distance, where $H_0 = f''(\theta_0)$ and Q_n is the distribution of $\sqrt{n}(\theta - \theta_n)$ when $\theta \sim \pi_n$. This strengthens previous BvM results for CL-posteriors by showing almost sure convergence (rather than convergence in probability) with respect to total variation distance (rather than in the weak topology).

10

To use Theorem 4.1, we also need to know the asymptotic distribution of θ_n. The asymptotics of θ_n are well-known (Lindsay, 1988; Varin et al., 2011), but for completeness we provide an informal derivation (see below). Define $G_j^n = \nabla_\theta\big|_{\theta=\theta_0} \log p_\theta^n(S_j^n|T_j^n)$. It turns out that $-\sqrt{n}(\theta_n - \theta_0) \approx \mathcal{N}(0, A_n^{-1}J_nA_n^{-1})$ under regularity conditions, where $A_n = \frac{1}{n}\sum_{j=1}^{k_n}\mathrm{Cov}(G_j^n)$ and $J_n = \frac{1}{n}\mathrm{Cov}\big(\sum_{j=1}^{k_n}G_j^n\big)$. Typically, $A_n \to H_0$ and $J_n \to J_0$ for some J_0, so that

$$-\sqrt{n}(\theta_n - \theta_0) \xrightarrow{\mathrm{D}} \mathcal{N}(0, H_0^{-1}J_0H_0^{-1}).$$

Hence, under typical conditions, the asymptotic distribution of $-\sqrt{n}(\theta_n - \theta_0)$ and the limit of Q_n are the same if and only if $H_0 = J_0$. Therefore, under these conditions, if $H_0 = J_0$ then the CL-posterior π_n has correct asymptotic frequentist coverage by Theorem 4.1. For instance, if for each n, $G_1^n, \ldots, G_{k_n}^n$ are pairwise uncorrelated, then $A_n = J_n$ and hence $H_0 = J_0$. However, in many cases of interest, $H_0 \neq J_0$ and the CL-posterior needs to be affinely transformed to have correct coverage (Ribatet et al., 2012; Pauli et al., 2011; Friel, 2012; Stoehr and Friel, 2015); also see Williams and Savitsky (2018) for a similar technique in survey sampling.

For completeness, here we provide a rough sketch of the derivation of the asymptotic distribution of θ_n; see Lindsay (1988) and Varin et al. (2011). By a first-order Taylor approximation applied to each entry of $f_n'(\theta) \in \mathbb{R}^D$, when θ_n is near θ_0 we have $0 = f_n'(\theta_n) \approx f_n'(\theta_0) + f_n''(\theta_0)(\theta_n - \theta_0)$, and thus, $-\sqrt{n}(\theta_n - \theta_0) \approx f_n''(\theta_0)^{-1}(\sqrt{n}f_n'(\theta_0))$, assuming $f_n''(\theta_0) \in \mathbb{R}^{D\times D}$ exists and is invertible and the error terms are negligible. When n is large, we typically have $f_n''(\theta_0) \approx \mathrm{E}f_n''(\theta_0)$ (for instance, due to a law of large numbers result), and thus, $f_n''(\theta_0) \approx \mathrm{E}f_n''(\theta_0) = \frac{1}{n}\sum_{j=1}^{k_n}\mathrm{E}(G_j^nG_j^{n\mathsf{T}}) = \frac{1}{n}\sum_{j=1}^{k_n}\mathrm{Cov}(G_j^n) = A_n$ since $\mathrm{E}(G_j^n) = 0$ and $\mathrm{E}\big(\nabla_\theta^2\big|_{\theta=\theta_0}\log p_\theta^n(S_j^n|T_j^n)\big) = -\mathrm{E}(G_j^nG_j^{n\mathsf{T}})$, as long as we can interchange the order of integrals and derivatives. Further, assuming a central limit theorem holds, $\sqrt{n}f_n'(\theta_0) = -\frac{1}{\sqrt{n}}\sum_{j=1}^{k_n}G_j^n \approx \mathcal{N}(0, J_n)$ where $J_n = \frac{1}{n}\mathrm{Cov}\big(\sum_{j=1}^{k_n}G_j^n\big)$. Thus, under appropriate conditions, $-\sqrt{n}(\theta_n - \theta_0) \approx \mathcal{N}(0, A_n^{-1}J_nA_n^{-1})$.

6 Applications to standard posteriors

In this section, we illustrate how our results can be used to easily prove posterior concentration, the Laplace approximation, and asymptotic normality for standard models such as exponential families, linear regression, and generalized linear models such as logistic regression and Poisson regression. We do not assume that the model is correctly specified.

6.1 Exponential families

Consider an exponential family with density $q(y|\eta) = \exp(\eta^\mathsf{T}s(y) - \kappa(\eta))$ with respect to a sigma-finite Borel measure λ on $\mathcal{Y} \subseteq \mathbb{R}^d$ where $s : \mathcal{Y} \to \mathbb{R}^k$, $\eta \in \mathcal{E} \subseteq \mathbb{R}^k$, and $\kappa(\eta) = \log\int_\mathcal{Y}\exp(\eta^\mathsf{T}s(y))\lambda(dy)$. Any exponential family on $\mathbb{R}^d$ can be put in this form by choosing λ appropriately and possibly reparametrizing to η. Let $Q_\eta(E) = \int_E q(y|\eta)\lambda(dy)$ and denote $\mathrm{E}_\eta s(Y) = \int_\mathcal{Y} s(y)Q_\eta(dy)$. For any $m \in \mathbb{N}$, we give $\mathbb{R}^m$ the Euclidean metric and the resulting Borel sigma-algebra unless otherwise specified.

Condition 6.1. *Assume $q(y|\eta)$ is of the form above, $\mathcal{E} = \{\eta \in \mathbb{R}^k : |\kappa(\eta)| < \infty\}$, $\mathcal{E}$ is open, $\mathcal{E}$ is nonempty, and $\eta \mapsto Q_\eta$ is one-to-one (that is, η is identifiable).*

Theorem 6.2 (Exponential families). *Consider a family $q(y|\eta)$ satisfying Condition 6.1. Suppose $Y_1, Y_2, \ldots \in \mathcal{Y}$ are i.i.d. random vectors such that $\mathrm{E}s(Y_i) = \mathrm{E}_{\theta_0}s(Y)$ for some $\theta_0 \in \Theta := \mathcal{E}$. Then for any open ball E such that $\theta_0 \in E$ and $\bar{E} \subseteq \Theta$, $f_n(\theta) := -\frac{1}{n}\sum_{i=1}^n \log q(Y_i|\theta)$ satisfies the conditions of Theorem 3.2 with probability 1.*

Proof. Note that $f_n(\theta) = \kappa(\theta) - \theta^\mathsf{T}S_n$ where $S_n = \frac{1}{n}\sum_{i=1}^n s(Y_i)$. By standard exponential family theory (e.g., Miller and Harrison, 2014, Prop. 19), κ is C^∞ (that is, κ has continuous derivatives of all order), κ is convex on Θ, $\kappa'(\theta) = \mathrm{E}_\theta s(Y)$, and $\kappa''(\theta)$ is symmetric positive definite for all $\theta \in \Theta$. Let $s_0 = \mathrm{E}s(Y_i)$. Since $s_0 = \mathrm{E}s(Y_i) = \mathrm{E}_{\theta_0}s(Y) = \kappa'(\theta_0)$ and $\kappa'(\theta_0)$ is finite (because κ is C^∞), $S_n \to s_0$ with probability 1 by the strong law of large numbers. Thus, letting $f(\theta) = \kappa(\theta) - \theta^\mathsf{T}s_0$, we have that with probability 1, for all $\theta \in \Theta$, $f_n(\theta) = \kappa(\theta) - \theta^\mathsf{T}S_n \to \kappa(\theta) - \theta^\mathsf{T}s_0 = f(\theta)$. (Note that due to the almost sure convergence of the sufficient statistics, we not only have that for all θ, with probability 1, $f_n(\theta) \to f(\theta)$, but we have the stronger consequence that with probability 1, for all θ, $f_n(\theta) \to f(\theta)$, which is needed for Theorem 3.2 to apply.) Let E be an open ball such that $\theta_0 \in E$ and $\bar{E} \subseteq \Theta$. Then $\kappa'''(\theta)$ is bounded on $\bar{E}$, since $\kappa'''(\theta)$ is continuous and $\bar{E}$ is compact. Hence, (f_n''') is uniformly bounded on E because $f_n'''(\theta) = \kappa'''(\theta)$. Therefore, with probability 1, $f_n \to f$ pointwise, f_n is convex and has continuous third derivatives on Θ, $f'(\theta_0) = \kappa'(\theta_0) - s_0 = 0$, $f''(\theta_0) = \kappa''(\theta_0)$ is positive definite, and (f_n''') is uniformly bounded on E. □

6.2 Generalized linear models (GLMs)

First, we state a general theorem for GLMs, then we show how it applies to commonly used GLMs. Consider a regression model of the form $p(y_i \mid \theta, x_i) \propto_\theta q(y_i \mid \theta^\mathsf{T}x_i)$ for covariates $x_i \in \mathcal{X} \subseteq \mathbb{R}^D$ and coefficients $\theta \in \Theta \subseteq \mathbb{R}^D$, where $q(y|\eta) = \exp(\eta s(y) - \kappa(\eta))$ is a one-parameter exponential family satisfying Condition 6.1. Note that the proportionality here is with respect to θ, not y_i. Assume Θ is open, Θ is convex, and $\theta^\mathsf{T}x \in \mathcal{E}$ for all $\theta \in \Theta$, $x \in \mathcal{X}$.

Theorem 6.3 (GLMs). *Suppose $(X_1, Y_1), (X_2, Y_2), \ldots \in \mathcal{X} \times \mathcal{Y}$ i.i.d. such that:*

1. *$f'(\theta_0) = 0$ for some $\theta_0 \in \Theta$, where $f(\theta) = -\mathrm{E}\log q(Y_i \mid \theta^\mathsf{T}X_i)$,*

2. *$\mathrm{E}|X_i s(Y_i)| < \infty$ and $\mathrm{E}|\kappa(\theta^\mathsf{T}X_i)| < \infty$ for all $\theta \in \Theta$,*

3. *for all $a \in \mathbb{R}^D$, if $a^\mathsf{T}X_i \overset{\mathrm{a.s.}}{=} 0$ then $a = 0$, and*

4. *there is an open ball $E \subseteq \mathbb{R}^D$ such that $\theta_0 \in E$, $\bar{E} \subseteq \Theta$, and for all $j, k, \ell \in \{1, \ldots, D\}$, $\mathrm{E}\big(\sup_{\theta \in \bar{E}} |\kappa'''(\theta^\mathsf{T}X_i)X_{ij}X_{ik}X_{i\ell}|\big) < \infty$.*

Then for any E as in assumption 4, $f_n(\theta) := -\frac{1}{n}\sum_{i=1}^n \log q(Y_i \mid \theta^\mathsf{T}X_i)$ satisfies the conditions of Theorem 3.2 with probability 1.

Proof. For all $\theta \in \Theta$, $f_n(\theta) = \frac{1}{n}\sum_{i=1}^n \kappa(\theta^\mathsf{T}X_i) - \theta^\mathsf{T}S_n$ where $S_n = \frac{1}{n}\sum_{i=1}^n X_i s(Y_i)$. Thus, $f_n(\theta)$ is C^∞ on Θ by the chain rule, since $\kappa(\eta)$ is C^∞ on $\mathcal{E}$. Further, $f_n(\theta)$ is convex since $\kappa(\eta)$ is convex. Noting that

$$f(\theta) = -\mathrm{E}\log q(Y_i \mid \theta^\mathsf{T}X_i) = \mathrm{E}(\kappa(\theta^\mathsf{T}X_i)) - \theta^\mathsf{T}\mathrm{E}(X_i s(Y_i)),$$

12

the assumed moment conditions (2) ensure that for all $\theta \in \Theta$, with probability 1, $f_n(\theta) \to f(\theta)$. This implies that with probability 1, for all $\theta \in \Theta$, $f_n(\theta) \to f(\theta)$, by the following argument. For any countable set $C \subseteq \Theta$, we have that with probability 1, for all $\theta \in C$, $f_n(\theta) \to f(\theta)$. Hence, letting C be a countable dense subset of Θ, and using the fact that each f_n is convex, we have that with probability 1, the limit $\tilde{f}(\theta) := \lim_n f_n(\theta)$ exists and is finite for all $\theta \in \Theta$ and $\tilde{f}$ is convex (Rockafellar, 1970, Thm 10.8). Since f is also convex, then $\tilde{f}$ and f are continuous functions (Rockafellar, 1970, Thm 10.1) that agree on a dense subset, so they are equal.

Choose E according to assumption 4. We show that with probability 1, (f_n''') is uniformly bounded on E. Fix $j, k, \ell \in \{1, \ldots, D\}$, and define $T(\theta, x) = \kappa'''(\theta^{\mathsf{T}}x)x_j x_k x_\ell$ for $\theta \in \Theta$, $x \in \mathcal{X}$. For all $x \in \mathcal{X}$, $\theta \mapsto T(\theta, x)$ is continuous, and for all $\theta \in \Theta$, $x \mapsto T(\theta, x)$ is measurable. Since $f_n'''(\theta)_{jk\ell} = \frac{1}{n}\sum_{i=1}^{n} T(\theta, X_i)$, assumption 4 implies that with probability 1, $(f_n'''(\theta)_{jk\ell})$ is uniformly bounded on $\bar{E}$, by the uniform law of large numbers (Ghosh and Ramamoorthi, 2003, Thm 1.3.3). Letting $C_{jk\ell}(X_1, X_2, \ldots)$ be such a uniform bound for each j, k, ℓ, we have that with probability 1, for all $n \in \mathbb{N}$, $\theta \in \bar{E}$, $\|f_n'''(\theta)\|^2 = \sum_{j,k,\ell} f_n'''(\theta)_{jk\ell}^2 \leq \sum_{j,k,\ell} C_{jk\ell}(X_1, X_2, \ldots)^2 < \infty$. Thus, (f_n''') is a.s. uniformly bounded on $\bar{E}$, and hence on E.

By Theorem 3.4, $f''(\theta_0) \overset{\text{a.s.}}{=} \lim_{n \to \infty} f_n''(\theta_0) = \lim \frac{1}{n}\sum_{i=1}^{n} \kappa''(\theta_0^{\mathsf{T}}X_i)X_i X_i^{\mathsf{T}}$. Since this limit exists and is finite almost surely, then by the strong law of large numbers, the limit must be equal to the expectation (Kallenberg, 2002, 4.23), that is, $f''(\theta_0) = \mathrm{E}\big(\kappa''(\theta_0^{\mathsf{T}}X_i)X_i X_i^{\mathsf{T}}\big)$. Thus, $f''(\theta_0)$ is positive definite, since for all nonzero $a \in \mathbb{R}^D$, $a^{\mathsf{T}}f''(\theta_0)a = \mathrm{E}\big(\kappa''(\theta_0^{\mathsf{T}}X_i)a^{\mathsf{T}}X_i X_i^{\mathsf{T}}a\big) > 0$, by the fact that $\kappa''(\eta) > 0$ for all $\eta \in \mathcal{E}$ and by assumption 3, $a^{\mathsf{T}}X_i X_i^{\mathsf{T}}a = |a^{\mathsf{T}}X_i|^2$ is strictly positive with positive probability. $\qquad\square$

6.2.1 Linear regression

The linear regression model is $p(y_i \mid \theta, x_i) = \mathcal{N}(y_i \mid \theta^{\mathsf{T}}x_i, \sigma^2)$ for $y_i \in \mathcal{Y} := \mathbb{R}$, $x_i \in \mathcal{X} := \mathbb{R}^D$, and $\theta \in \Theta := \mathbb{R}^D$. The model can equivalently be written as $p(y_i \mid \theta, x_i) \propto_\theta q(y_i \mid \theta^{\mathsf{T}}x_i)$ where $q(y|\eta) := \exp(\eta s(y) - \kappa(\eta))$ is a density with respect to $\lambda(dy) = \mathcal{N}(y \mid 0, \sigma^2)dy$ for $y \in \mathcal{Y}$ and $\eta \in \mathcal{E} := \mathbb{R}$, by defining $s(y) = y/\sigma^2$ and $\kappa(\eta) = \eta^2/(2\sigma^2)$.

Theorem 6.4 (Linear regression). *Suppose $(X_1, Y_1), (X_2, Y_2), \ldots \in \mathcal{X} \times \mathcal{Y}$ i.i.d. such that:*

1. *$\mathrm{E}|X_iY_i| < \infty$, $\mathrm{E}\|X_i X_i^{\mathsf{T}}\| < \infty$, and*

2. *for all $a \in \mathbb{R}^D$, if $a^{\mathsf{T}}X_i \overset{\text{a.s.}}{=} 0$ then $a = 0$.*

Then $\theta_0 := (\mathrm{E}X_i X_i^{\mathsf{T}})^{-1}\mathrm{E}X_iY_i$ is well-defined and for any open ball E such that $\theta_0 \in E$, $f_n(\theta) := -\frac{1}{n}\sum_{i=1}^{n} \log q(Y_i \mid \theta^{\mathsf{T}}X_i)$ satisfies the conditions of Theorem 3.2 with probability 1.

Proof. For any random vector $Z \in \mathbb{R}^k$, $\mathrm{E}|Z| < \infty$ if and only if $\mathrm{E}Z$ exists and is finite; likewise for matrices and tensors. Thus, $\mathrm{E}X_iY_i$ and $\mathrm{E}X_i X_i^{\mathsf{T}}$ exist and are finite. Further, $\mathrm{E}X_i X_i^{\mathsf{T}}$ is positive definite (and hence, invertible) since for all nonzero $a \in \mathbb{R}^D$, $a^{\mathsf{T}}(\mathrm{E}X_i X_i^{\mathsf{T}})a = \mathrm{E}|a^{\mathsf{T}}X_i|^2 > 0$. Condition 6.1 is easily checked: $\mathcal{E} = \{\eta \in \mathbb{R} : |\kappa(\eta)| < \infty\}$ since $\eta^2/(2\sigma^2) < \infty$ for all $\eta \in \mathbb{R}$, $\mathcal{E}$ is open and nonempty, and the mean of a normal distribution is identifiable. The GLM conditions are also straightforward to verify. Θ is open and convex, and $\theta^{\mathsf{T}}x \in \mathcal{E}$ for all $\theta \in \Theta$, $x \in \mathcal{X}$. Condition 3 of Theorem 6.3 is satisfied by assumption, and condition 4 is satisfied trivially since $\kappa'''(\eta) = 0$ for all $\eta \in \mathcal{E}$. Assumption 1 implies that

13

condition 2 of Theorem 6.3 holds, since $\mathrm{E}|X_i s(Y_i)| = \mathrm{E}|X_i Y_i|/\sigma^2 < \infty$ and $\mathrm{E}|\kappa(\theta^\mathsf{T} X_i)| = \theta^\mathsf{T}(\mathrm{E}X_i X_i^\mathsf{T})\theta/(2\sigma^2) < \infty$. Finally, it is straightforward to verify that condition 1 of Theorem 6.3 holds with $\theta_0 = (\mathrm{E}X_i X_i^\mathsf{T})^{-1}\mathrm{E}X_i Y_i$. □

6.2.2 Logistic regression

The logistic regression model is $p(y_i \mid \theta, x_i) = \mathrm{Bernoulli}(y_i \mid \sigma(\theta^\mathsf{T} x_i))$ for $y_i \in \mathcal{Y} := \{0, 1\}$, $x_i \in \mathcal{X} := \mathbb{R}^D$, and $\theta \in \Theta := \mathbb{R}^D$, where $\sigma(\eta) = 1/(1 + e^{-\eta})$ for $\eta \in \mathcal{E} := \mathbb{R}$. Thus, $p(y_i \mid \theta, x_i) = q(y_i \mid \theta^\mathsf{T} x_i)$ where $q(y|\eta) := \exp(\eta y - \kappa(\eta))$ is a density with respect to $\lambda = \delta_0 + \delta_1$ for $y \in \mathcal{Y}$ and $\eta \in \mathcal{E}$, by defining $\kappa(\eta) = \log(1 + e^\eta)$. Here, δ_y denotes the unit point mass at y.

Theorem 6.5 (Logistic regression). *Suppose* $(X_1, Y_1), (X_2, Y_2), \ldots \in \mathcal{X} \times \mathcal{Y}$ *i.i.d. such that:*

1. $f'(\theta_0) = 0$ *for some* $\theta_0 \in \Theta$, *where* $f(\theta) = -\mathrm{E}\log q(Y_i \mid \theta^\mathsf{T} X_i)$,

2. $\mathrm{E}|X_{ij}X_{ik}X_{i\ell}| < \infty$ *for all* $j, k, \ell \in \{1, \ldots, D\}$, *and*

3. *for all* $a \in \mathbb{R}^D$, *if* $a^\mathsf{T} X_i \overset{\text{a.s.}}{=} 0$ *then* $a = 0$.

Then for any open ball $E \subseteq \Theta$ *such that* $\theta_0 \in E$, $f_n(\theta) := -\frac{1}{n}\sum_{i=1}^n \log q(Y_i \mid \theta^\mathsf{T} X_i)$ *satisfies the conditions of Theorem 3.2 with probability 1.*

Proof. Condition 6.1 is easily checked: $\mathcal{E} = \{\eta \in \mathbb{R} : |\kappa(\eta)| < \infty\}$, $\mathcal{E}$ is open and nonempty, and η is identifiable since $\sigma(\eta)$ is one-to-one. Trivially, Θ is open and convex, and $\theta^\mathsf{T} x \in \mathcal{E}$ for all $\theta \in \Theta$, $x \in \mathcal{X}$. Conditions 1 and 3 of Theorem 6.3 are satisfied by assumptions 1 and 3, respectively. Condition 4 of Theorem 6.3 is satisfied due to assumption 2 and the fact that $|\kappa'''(\eta)| \leq 3$ for all $\eta \in \mathcal{E}$, because $\kappa''' = \sigma(1 - \sigma)(1 - 2\sigma)^2 - 2\sigma^2(1 - \sigma)^2$ and $0 < \sigma(\eta) < 1$. Assumption 2 also implies that $\mathrm{E}|X_i| < \infty$, because $|X_i| \leq \sum_j |X_{ij}|$ and $\mathrm{E}|X_{ij}| < \infty$ for all j (Folland, 2013, 6.12). It follows that condition 2 of Theorem 6.3 holds, since $\mathrm{E}|X_i Y_i| \leq \mathrm{E}|X_i| < \infty$ and $\mathrm{E}|\kappa(\theta^\mathsf{T} X_i)| \leq \log 2 + \mathrm{E}|\theta^\mathsf{T} X_i| \leq \log 2 + |\theta|\mathrm{E}|X_i| < \infty$, where we have used the inequality $|\kappa(\eta)| = \log(1 + e^\eta) \leq \log 2 + |\eta|$ for $\eta \in \mathbb{R}$. □

6.2.3 Poisson regression

The Poisson regression model is $p(y_i \mid \theta, x_i) = \mathrm{Poisson}(y_i \mid \exp(\theta^\mathsf{T} x_i))$ for $y_i \in \mathcal{Y} := \{0, 1, 2, \ldots\}$, $x_i \in \mathcal{X} := \mathbb{R}^D$, and $\theta \in \Theta := \mathbb{R}^D$. Thus, $p(y_i \mid \theta, x_i) \propto_\theta q(y_i \mid \theta^\mathsf{T} x_i)$ where $q(y|\eta) := \exp(\eta y - \kappa(\eta))$ is a density with respect to $\lambda := \sum_{y \in \mathcal{Y}} \delta_y/y!$ for $y \in \mathcal{Y}$ and $\eta \in \mathcal{E} := \mathbb{R}$, by defining $\kappa(\eta) = e^\eta$.

Theorem 6.6 (Poisson regression). *Suppose* $(X_1, Y_1), (X_2, Y_2), \ldots \in \mathcal{X} \times \mathcal{Y}$ *i.i.d. such that:*

1. $f'(\theta_0) = 0$ *for some* $\theta_0 \in \Theta$, *where* $f(\theta) = -\mathrm{E}\log q(Y_i \mid \theta^\mathsf{T} X_i)$,

2. $\mathrm{E}|X_i Y_i| < \infty$ *and* $\mathrm{E}\exp(c|X_i|) < \infty$ *for all* $c > 0$, *and*

3. *for all* $a \in \mathbb{R}^D$, *if* $a^\mathsf{T} X_i \overset{\text{a.s.}}{=} 0$ *then* $a = 0$.

Then for any open ball $E \subseteq \Theta$ *such that* $\theta_0 \in E$, $f_n(\theta) := -\frac{1}{n}\sum_{i=1}^n \log q(Y_i \mid \theta^\mathsf{T} X_i)$ *satisfies the conditions of Theorem 3.2 with probability 1.*

14

Proof. As before, Condition 6.1 is easily checked: $\mathcal{E} = \{\eta \in \mathbb{R} : |\kappa(\eta)| < \infty\}$, $\mathcal{E}$ is open and nonempty, and η is identifiable. Trivially, Θ is open and convex, and $\theta^{\mathsf{T}}x \in \mathcal{E}$ for all $\theta \in \Theta$, $x \in \mathcal{X}$. Conditions 1 and 3 of Theorem 6.3 are satisfied by assumptions 1 and 3. Condition 2 of Theorem 6.3 is satisfied due to assumption 2, since for all $\theta \in \Theta$, $\mathrm{E}|\kappa(\theta^{\mathsf{T}}X_i)| = \mathrm{E}\exp(\theta^{\mathsf{T}}X_i) \leq \mathrm{E}\exp(|\theta||X_i|) < \infty$. For all $m \in \mathbb{N}$ and $j \in \{1,\ldots,D\}$, $\mathrm{E}|X_{ij}|^m \leq \mathrm{E}|X_i|^m = m!\mathrm{E}(|X_i|^m/m!) \leq m!\mathrm{E}\exp(|X_i|) < \infty$. Further, letting $r > 0$, $c = |\theta_0| + r$, and $E = B_r(\theta_0)$, we have that for all $\theta \in \bar{E}$, $\kappa'''(\theta^{\mathsf{T}}X_i) = \exp(\theta^{\mathsf{T}}X_i) \leq \exp(c|X_i|)$. Hence,

$$\mathrm{E}\big(\sup_{\theta \in \bar{E}} |\kappa'''(\theta^{\mathsf{T}}X_i)X_{ij}X_{ik}X_{i\ell}|\big) \leq \mathrm{E}\big(e^{c|X_i|}|X_{ij}X_{ik}X_{i\ell}|\big) \leq \Big(\mathrm{E}e^{4c|X_i|}\mathrm{E}|X_{ij}|^4\mathrm{E}|X_{ik}|^4\mathrm{E}|X_{i\ell}|^4\Big)^{1/4}$$

by Hölder's inequality (Folland, 2013, 6.2); thus, condition 4 of Theorem 6.3 is satisfied. $\square$

7 Applications to generalized posteriors

7.1 Pseudolikelihood-based posteriors

Pseudolikelihood (Besag, 1975) is a powerful approach for many models in which the likelihood is difficult to compute due to intractability of the normalization constant. Instead of the standard likelihood $L(\theta) = p(y_1,\ldots,y_n \mid \theta)$, the basic idea is to use a *pseudolikelihood* $\mathcal{L}(\theta) = \prod_{i=1}^n p(y_i \mid y_{-i},\theta)$ where $y_{-i} = (y_1,\ldots,y_{i-1},y_{i+1},\ldots,y_n)$. Maximum pseudolikelihood estimates are used in many applications and have been shown to be consistent and asymptotically normal in a range of cases (Besag, 1975; Geman and Graffigne, 1986; Gidas, 1988; Comets, 1992; Jensen and Künsch, 1994; Mase, 1995; Liang and Yu, 2003; Hyvärinen, 2006). Usage of pseudolikelihoods for constructing generalized posteriors is much less common, perhaps due to concerns about the validity of the resulting posterior (but see Zhou and Schmidler, 2009; Bouranis et al., 2017; Pauli et al., 2011; Rydén and Titterington, 1998).

In this section, we provide sufficient conditions for concentration, asymptotic normality, and the Laplace approximation for a large class of pseudolikelihood-based posteriors. Specifically, we consider pseudolikelihoods in which each factor takes the form of a generalized linear model. We provide a general result for pseudolikelihoods in this class, and then consider three cases in particular: the Ising model on $\mathbb{Z}^m$ (Section 7.2), Gaussian Markov random fields (Section 7.3), and fully visible Boltzmann machines (Section 7.4). Since any pseudolikelihood is a composite likelihood, the consistency and coverage results discussed in Section 5 apply here.

Condition 7.1. *Suppose the data can be arranged in a sequence $y_1, y_2, \ldots \in \mathcal{Y} \subseteq \mathbb{R}^d$ and consider a pseudolikelihood of the form:*

$$\mathcal{L}_n^{\mathrm{pseudo}}(\theta) \propto \prod_{i=1}^n q\big(y_i \mid \theta^{\mathsf{T}}\varphi_i(\vec{y})\big)$$

for $\theta \in \Theta \subseteq \mathbb{R}^D$, where $\varphi_i(\vec{y}) \in \mathcal{X} \subseteq \mathbb{R}^D$ is a function of $\vec{y} = (y_1, y_2, \ldots)$ and $q(y|\eta) = \exp(\eta s(y) - \kappa(\eta))$ is a one-parameter exponential family satisfying Condition 6.1 for $y \in \mathcal{Y}$, $\eta \in \mathcal{E}$. Assume Θ is open and convex, and $\theta^{\mathsf{T}}x \in \mathcal{E}$ for all $\theta \in \Theta$, $x \in \mathcal{X}$.

15

Theorem 7.2. *Assume the setup in Condition 7.1. Let $\vec{Y} = (Y_1, Y_2, \ldots)$ be a sequence of random vectors in $\mathcal{Y}$ and define $X_i = \varphi_i(\vec{Y})$. Suppose $(X_1, Y_1), (X_2, Y_2), \ldots$ are identically distributed, but not necessarily independent. Define $f_n(\theta) = -\frac{1}{n} \sum_{i=1}^{n} \log q(Y_i \mid \theta^{\mathsf{T}} X_i)$ and $f(\theta) = -\mathrm{E} \log q(Y_i \mid \theta^{\mathsf{T}} X_i)$ for $\theta \in \Theta$. Assume:*

1. *for all $\theta \in \Theta$, $f(\theta)$ is finite and $f_n(\theta) \xrightarrow{\text{a.s.}} f(\theta)$ as $n \to \infty$,*

2. *there exists $\theta_0 \in \Theta$ such that $f'(\theta_0) = 0$ and $f''(\theta_0) = \mathrm{E}\big(\kappa''(\theta_0^{\mathsf{T}} X_i) X_i X_i^{\mathsf{T}}\big)$,*

3. *for all $a \in \mathbb{R}^D$, if $a^{\mathsf{T}} X_i \overset{\text{a.s.}}{=} 0$ then $a = 0$, and*

4. *with probability 1, (f_n''') is uniformly bounded on some open ball $E \subseteq \Theta$ containing θ_0.*

Then for any E as in assumption 4, f_n satisfies the conditions of Theorem 3.2 with probability 1.

Remark. To explain the notation, we assume the observed data consist of the first n elements of single random sequence $\vec{Y} = (Y_1, Y_2, \ldots)$, where each Y_i is a random vector. In the Ising model and Gaussian Markov random field (GMRF) examples (Sections 7.2 and 7.3), $\vec{Y}$ contains the values at the nodes of a single infinite graph, arranged as a sequence. In the GMRF example, note that this infinite graph could consist of a single connected component (as in the Ising model) or could consist of many disconnected components, for instance, many samples of a single finite or infinite graph. In both examples, the assumption that $(X_1, Y_1), (X_2, Y_2), \ldots$ are identically distributed (but not necessarily independent) is similar to a stationarity condition, but is weaker than stationarity. This condition is fairly mild since stationarity is commonly assumed. In the fully visible Boltzmann machine (Section 7.4), we assume i.i.d. samples of graphs, which simplifies the analysis in some respects.

Proof. As in the proof of Theorem 6.3, f_n is C^∞, f_n is convex, and by convexity, assumption 1 implies that with probability 1, for all $\theta \in \Theta$, $f_n(\theta) \to f(\theta)$. By Theorem 3.4, $f''(\theta_0)$ exists and is finite. Thus, $f''(\theta_0)$ is positive definite since for all nonzero $a \in \mathbb{R}^D$, $a^{\mathsf{T}} f''(\theta_0) a = \mathrm{E}\big(\kappa''(\theta_0^{\mathsf{T}} X_i) a^{\mathsf{T}} X_i X_i^{\mathsf{T}} a\big) > 0$ by assumptions 2 and 3 and the fact that $\kappa''(\eta) > 0$. $\qquad\square$

7.2 Ising model

The Ising model is a classical model of ferromagnetism in statistical mechanics and has gained widespread use in many other applications such as spatial statistics (Banerjee et al., 2014) and image processing (Geman and Geman, 1984). Pseudolikelihood-based posteriors for the Ising model and Potts model, more generally, have been used by Zhou and Schmidler (2009) for protein modeling.

Consider the m-dimensional integer lattice $\mathbb{Z}^m$ and let $v : \mathbb{N} \to \mathbb{Z}^m$ be a bijection from $\mathbb{N}$ to $\mathbb{Z}^m$. Let $y_1, y_2, \ldots \in \mathcal{Y} := \{-1, 1\}$ be variables associated with the points of $\mathbb{Z}^m$ such that y_i is the value at $v(i)$. The Ising model is a Markov random field with singleton potentials $\exp(\theta_1 y_i)$ for each $i \in \mathbb{N}$ and pairwise potentials $\exp(\theta_2 y_i y_j)$ for each pair $i, j \in \mathbb{N}$ such that $v(i)$ and $v(j)$ are adjacent in $\mathbb{Z}^m$, that is, such that $|v(i) - v(j)| = 1$. This motivates the use of the pseudolikelihood (Besag, 1975),

$$\mathcal{L}_n^{\text{Ising}}(\theta) = \prod_{i=1}^{n} \frac{\exp(\theta_1 y_i + \theta_2 \sum_{j \in N_i} y_i y_j)}{\sum_{y \in \mathcal{Y}} \exp(\theta_1 y + \theta_2 \sum_{j \in N_i} y y_j)}$$

16

for $\theta \in \Theta := \mathbb{R}^2$, where $N_i = \{j \in \mathbb{N} : v(j) \text{ is adjacent to } v(i)\}$. By defining $q(y|\eta) = \exp(\eta y - \kappa(\eta))$ for $y \in \{-1, 1\}$ and $\eta \in \mathbb{R}$, where $\kappa(\eta) = \log(e^\eta + e^{-\eta})$, the Ising model pseudolikelihood can be written as $\mathcal{L}_n^{\text{Ising}}(\theta) = \prod_{i=1}^n q(y_i \mid \theta_1 + \theta_2 \sum_{j \in N_i} y_j)$.

Theorem 7.3. *Let $\vec{Y} = (Y_1, Y_2, \ldots)$ be a sequence of random variables in $\{-1, 1\}$ and define $X_i = \left(1, \sum_{j \in N_i} Y_j\right)^{\mathsf{T}} \in \mathbb{R}^2$. Suppose $(X_1, Y_1), (X_2, Y_2), \ldots$ are identically distributed, but not necessarily independent. Define $f_n(\theta) = -\frac{1}{n} \sum_{i=1}^n \log q(Y_i \mid \theta^{\mathsf{T}} X_i)$ and $f(\theta) = -\mathrm{E} \log q(Y_i \mid \theta^{\mathsf{T}} X_i)$ for $\theta \in \Theta$. Assume:*

   1. for all $\theta \in \Theta$, $f_n(\theta) \xrightarrow{\text{a.s.}} f(\theta)$ as $n \to \infty$,

   2. $f'(\theta_0) = 0$ for some $\theta_0 \in \Theta$, and

   3. $\mathrm{Var}\left(\sum_{j \in N_i} Y_j\right) > 0$.

Then for any open ball E such that $\theta_0 \in E$, f_n satisfies the conditions of Theorem 3.2 with probability 1.

Proof. We apply Theorem 7.2. Define $\mathcal{X} = \left\{(1, z)^{\mathsf{T}} : z \in \{-2m, \ldots, 2m\}\right\}$, noting that $X_i \in \mathcal{X}$. It is easy to check that Condition 7.1 holds. For all $\theta \in \Theta$, $f(\theta)$ is finite since $|\mathcal{X} \times \mathcal{Y}| < \infty$. If $a^{\mathsf{T}} X_i \overset{\text{a.s.}}{=} 0$ then $a = 0$, since $a^{\mathsf{T}} X_i = a_1 + a_2 \sum_{j \in N_i} Y_j$ and $\mathrm{Var}\left(\sum_{j \in N_i} Y_j\right) > 0$. Let E be an open ball containing θ_0, and let $c = \sup\{|\kappa'''(\theta^{\mathsf{T}} x)| : x \in \mathcal{X}, \theta \in \bar{E}\}$. Then $c < \infty$ since κ''' is continuous, $|\mathcal{X}|$ is finite, and $\bar{E}$ is compact. Therefore, for all $\theta \in E$, $|f_n'''(\theta)_{jk\ell}| \leq \frac{1}{n} \sum_{i=1}^n |\kappa'''(\theta^{\mathsf{T}} X_i) X_{ij} X_{ik} X_{i\ell}| \leq c(2m)^3$, and thus, (f_n''') is a.s. uniformly bounded on E. Finally, $f''(\theta_0) = \mathrm{E}\left(\kappa''(\theta_0^{\mathsf{T}} X_i) X_i X_i^{\mathsf{T}}\right)$ because differentiating under the integral sign is justified by the bounds $|\kappa(\eta)| \leq |\eta| + \log 2$, $|\kappa'(\eta)| \leq 1$, $|\kappa''(\eta)| \leq 2$, and $|X_{ij}| \leq 2m$ (Folland, 2013, 2.27). $\square$

7.3   Gaussian Markov random fields

Gaussian Markov random fields are widely used in spatial statistics and time-series (Banerjee et al., 2014). Let G be an infinite regular graph with vertices $v(1), v(2), \ldots$, and let $y_1, y_2, \ldots \in \mathbb{R}$ be variables associated with the vertices of G such that y_i is the value at $v(i)$. Consider a model in which the conditional distribution of y_i given y_{-i} is $p_\theta(y_i|y_{-i}) = \mathcal{N}(y_i \mid \theta^{\mathsf{T}} \varphi_i(\vec{y}), \gamma^{-1})$ where $\theta \in \Theta := \mathbb{R}^D$, $\varphi_i(\vec{y}) = (y_j : j \in N_i) \in \mathbb{R}^D$, and $N_i = \{j \in \mathbb{N} : v(j) \text{ is adjacent to } v(i)\}$. This leads to the pseudolikelihood (Besag, 1975)

$$\mathcal{L}_n^{\text{GRF}}(\theta) = \prod_{i=1}^n p_\theta(y_i|y_{-i}) = \prod_{i=1}^n \mathcal{N}\left(y_i \mid \theta^{\mathsf{T}} \varphi_i(\vec{y}), \gamma^{-1}\right).$$

By defining $q(y|\eta) = \exp(\eta \gamma y - \kappa(\eta))$ for $y \in \mathbb{R}$ and $\eta \in \mathbb{R}$, where $\kappa(\eta) = \frac{1}{2} \gamma \eta^2$, this pseudolikelihood can be written as $\mathcal{L}_n^{\text{GRF}}(\theta) \propto \prod_{i=1}^n q\left(y_i \mid \theta^{\mathsf{T}} \varphi_i(\vec{y})\right)$.

Theorem 7.4. *Let $\vec{Y} = (Y_1, Y_2, \ldots)$ be a sequence of random variables in $\mathbb{R}$ and define $X_i = (Y_j : j \in N_i) \in \mathbb{R}^D$. Suppose $(X_1, Y_1), (X_2, Y_2), \ldots$ are identically distributed, but not necessarily independent. Assume:*

17

1. $\frac{1}{n}\sum_{i=1}^{n} X_i Y_i \xrightarrow{\text{a.s.}} EX_i Y_i \in \mathbb{R}^D$ *and* $\frac{1}{n}\sum_{i=1}^{n} X_i X_i^{\mathsf{T}} \xrightarrow{\text{a.s.}} EX_i X_i^{\mathsf{T}} \in \mathbb{R}^{D\times D}$ *as* $n \to \infty$, *and*

2. *for all* $a \in \mathbb{R}^D$, *if* $a^{\mathsf{T}} X_i \overset{\text{a.s.}}{=} 0$ *then* $a = 0$.

Then $\theta_0 := (EX_i X_i^{\mathsf{T}})^{-1} EX_i Y_i$ *is well-defined and for any open ball* E *such that* $\theta_0 \in E$, $f_n(\theta) := -\frac{1}{n}\sum_{i=1}^{n} \log q(Y_i \mid \theta^{\mathsf{T}} X_i)$ *satisfies the conditions of Theorem 3.2 with probability* 1.

Proof. We apply Theorem 7.2. Let $f(\theta) = -E\log q(Y_i \mid \theta^{\mathsf{T}} X_i) = \frac{1}{2}\gamma\theta^{\mathsf{T}}(EX_i X_i^{\mathsf{T}})\theta - \gamma\theta^{\mathsf{T}} EX_i Y_i$ for $\theta \in \mathbb{R}^D$. Thus, $f''(\theta) = \gamma(EX_i X_i^{\mathsf{T}}) = E(\kappa''(\theta^{\mathsf{T}} X_i)X_i X_i^{\mathsf{T}})$ since $\kappa''(\eta) = \gamma$. By assumption 1, for all $\theta \in \mathbb{R}^D$, $f(\theta)$ is finite and $f_n(\theta) \xrightarrow{\text{a.s.}} f(\theta)$ as $n \to \infty$. As in the case of linear regression (Theorem 6.4), $EX_i X_i^{\mathsf{T}}$ is positive definite by assumption 2, $f'(\theta_0) = 0$, and (f_n''') is a.s. uniformly bounded on all of $\mathbb{R}^D$ since $\kappa'''(\eta) = 0$. $\qquad\square$

7.4 Fully visible Boltzmann machines

The Boltzmann machine is a stochastic recurrent neural network originally developed as a model of neural computation (Hinton and Sejnowski, 1983; Ackley et al., 1985). Maximum pseudolikelihood estimation has been shown to be consistent for fully visible Boltzmann machines (Hyvärinen, 2006). Here, we consider the corresponding pseudolikelihood-based generalized posteriors.

Define $p_{A,b}(y) \propto \exp(y^{\mathsf{T}} Ay + b^{\mathsf{T}} y)$ for $y \in \mathcal{Y} := \{-1, 1\}^d$, where $A \in \mathbb{R}^{d\times d}$ is a strictly upper triangular matrix and $b \in \mathbb{R}^d$. Given samples from $p_{A,b}$, inference for A and b is complicated by the intractability of the normalization constant $Z_{A,b} = \sum_{y\in\mathcal{Y}} \exp(y^{\mathsf{T}} Ay + b^{\mathsf{T}} y)$ since $|\mathcal{Y}| = 2^d$ is very large when d is large. Observe that we can write

$$p_{A,b}(y_j|y_{-j}) \propto_{y_j} \exp\left(\sum_{k=1}^{j-1} A_{kj} y_k y_j + \sum_{k=j+1}^{d} A_{jk} y_j y_k + b_j y_j\right) = \exp\left(y_j \theta^{\mathsf{T}} \varphi_j(y)\right) \quad (6)$$

where $\theta = \theta(A, b) \in \mathbb{R}^D$ is a $D = d + d(d-1)/2$ dimensional vector concatenating b and the strictly upper triangular entries of A, and $\varphi_j(y) \in \{-1, 0, 1\}^D$ is a function that does not depend on y_j. Thus, we have $p_{A,b}(y_j|y_{-j}) = q(y_j \mid \theta^{\mathsf{T}} \varphi_j(y))$ by defining $q(y_j|\eta) = \exp(\eta y_j - \kappa(\eta))$ for $y_j \in \{-1, 1\}$ and $\eta \in \mathbb{R}$, where $\kappa(\eta) = \log(e^\eta + e^{-\eta})$. Now, suppose we have n samples $y_1, \ldots, y_n \in \mathcal{Y} = \{-1, 1\}^d$ and for $\theta \in \Theta := \mathbb{R}^D$, consider the pseudolikelihood

$$\mathcal{L}_n^{\text{Boltz}}(\theta) = \prod_{i=1}^{n}\prod_{j=1}^{d} p_{A,b}(y_{ij}|y_{i,-j}) = \prod_{i=1}^{n}\prod_{j=1}^{d} q(y_{ij} \mid \theta^{\mathsf{T}} \varphi_j(y_i)).$$

Theorem 7.5. *Let* $Y_1, Y_2, \ldots \in \mathcal{Y}$ *be i.i.d. random vectors and define* $X_{ij} = \varphi_j(Y_i)$. *Define* $f(\theta) = -\sum_{j=1}^{d} E\log q(Y_{ij} \mid \theta^{\mathsf{T}} X_{ij})$ *for* $\theta \in \Theta$. *Assume:*

1. $f'(\theta_0) = 0$ *for some* $\theta_0 \in \Theta$, *and*

2. *for all nonzero* $a \in \mathbb{R}^d$, $\text{Var}(a^{\mathsf{T}} Y_i) > 0$.

Then for any open ball E *such that* $\theta_0 \in E$, $f_n(\theta) := -\frac{1}{n}\sum_{i=1}^{n}\sum_{j=1}^{d} \log q(Y_{ij} \mid \theta^{\mathsf{T}} X_{ij})$ *satisfies the conditions of Theorem 3.2 with probability* 1.

Proof. Observe that $f_n(\theta) = \frac{1}{n}\sum_{i=1}^n \sum_{j=1}^d \kappa(\theta^\mathsf{T} X_{ij}) - \theta^\mathsf{T}\left(\frac{1}{n}\sum_{i=1}^n \sum_{j=1}^d X_{ij}Y_{ij}\right)$ and $f(\theta) = \sum_{j=1}^d \mathrm{E}\kappa(\theta^\mathsf{T} X_{ij}) - \theta^\mathsf{T}\left(\sum_{j=1}^d \mathrm{E}X_{ij}Y_{ij}\right)$. As in the proof of Theorem 6.3, f_n is C^∞ and convex. Since $\{-1,0,1\}^D$ is a finite set, $\sup\left\{|\kappa(\theta^\mathsf{T} x)| : x \in \{-1,0,1\}^D\right\} < \infty$ for all $\theta \in \Theta$. Also, $|X_{ijk}Y_{ij}| \leq 1$, and thus, $f(\theta)$ is finite and $f_n(\theta) \xrightarrow{\text{a.s.}} f(\theta)$ by the strong law of large numbers. (Note that the standard strong law of large numbers applies here since the data consist of n i.i.d. samples from the Boltzmann machine, rather than the first n elements of a single sample as in the Ising model and GMRF examples.) Due to convexity, this implies that with probability 1, for all $\theta \in \Theta$, $f_n(\theta) \to f(\theta)$ as $n \to \infty$.

Let E be an open ball containing θ_0. Then for all $\theta \in E$, $|f_n'''(\theta)_{k\ell m}| \leq cd$ where $c = \sup\{|\kappa'''(\theta^\mathsf{T} x)| : x \in \{-1,0,1\}^D, \theta \in \bar{E}\}$, and $c < \infty$ because κ''' is continuous and $\bar{E}$ is compact. Thus, for all $\theta \in E$, $\|f_n'''(\theta)\|^2 = \sum_{k,\ell,m} |f_n'''(\theta)_{k\ell m}|^2 \leq c^2 d^2 D^3$. Hence, (f_n''') is uniformly bounded on E.

Finally, we show that $f''(\theta_0)$ is positive definite. First, $f''(\theta_0) = \sum_{j=1}^d \mathrm{E}\left(\kappa''(\theta_0^\mathsf{T} X_{ij})X_{ij}X_{ij}^\mathsf{T}\right)$ because differentiating under the integral sign is justified by the bounds $|\kappa(\eta)| \leq |\eta| + \log 2$, $|\kappa'(\eta)| \leq 1$, $|\kappa''(\eta)| \leq 2$, and $|X_{ijk}| \leq 1$ (Folland, 2013, 2.27). Let $\theta \in \mathbb{R}^D$ be nonzero and let A, b be the corresponding parameters such that $\theta = \theta(A,b)$. Then by Equation 6, $A^\mathsf{T} Y_i + AY_i + b = (\theta^\mathsf{T} X_{i1}, \ldots, \theta^\mathsf{T} X_{id})^\mathsf{T} \in \mathbb{R}^d$. If $A \neq 0$, then $\mathrm{Var}(\theta^\mathsf{T} X_{ij'}) > 0$ for some j' by assumption 2, and hence, $\theta^\mathsf{T} f''(\theta_0)\theta = \sum_{j=1}^d \mathrm{E}\left(\kappa''(\theta_0^\mathsf{T} X_{ij})|\theta^\mathsf{T} X_{ij}|^2\right) > 0$ because $\kappa''(\eta) > 0$ and $\mathbb{P}(|\theta^\mathsf{T} X_{ij'}| > 0) > 0$. Meanwhile, if $A = 0$, then $b_{j'} \neq 0$ for some j' (because $\theta \neq 0$), and again $\theta^\mathsf{T} f''(\theta_0)\theta > 0$ because $|\theta^\mathsf{T} X_{ij'}| = |b_{j'}| > 0$. Therefore, $f''(\theta_0)$ is positive definite. $\square$

7.5   Cox proportional hazards model

The Cox proportional hazards model (Cox, 1972) is widely used for survival analysis. The proportional hazards model assumes the hazard function for subject i is $\lambda_0(y)\exp(\theta^\mathsf{T} x_i)$ for $y \geq 0$, where $\lambda_0(y) \geq 0$ is a baseline hazard function shared by all subjects, $x_i \in \mathbb{R}^D$ is a vector of covariates for subject i, and $\theta \in \mathbb{R}^D$ is a vector of coefficients. To perform inference for θ in a way that does not require any modeling of λ_0 and elegantly handles censoring, Cox (1972, 1975) proposed using the *partial likelihood*,

$$\mathcal{L}_n^{\mathrm{Cox}}(\theta) = \prod_{i=1}^n \left(\frac{\exp(\theta^\mathsf{T} x_i)}{\sum_{j=1}^n \exp(\theta^\mathsf{T} x_j)\mathbb{1}(y_j \geq y_i)}\right)^{z_i}$$

where $y_i \geq 0$ is the outcome time for subject i and $z_i \in \{0,1\}$ indicates whether y_i is an observed event time ($z_i = 1$) or a right-censoring time ($z_i = 0$). When $z_i = 1$, the ith factor in the partial likelihood can be interpreted as the conditional probability that subject i has an event at time y_i, given the risk set $\{j : y_j \geq y_i\}$ (the set of subjects that have not yet had an event or been censored up until time y_i) and given that some subject has an event at time y_i. See Efron (1977) for an intuitive explanation of the Cox partial likelihood based on a discrete approximation. Formally, the Cox partial likelihood coincides with the likelihood of a certain generalized linear model with categorical outcomes, however, asymptotic analysis is complicated by the dependencies between the factors of the partial likelihood. A number of authors have studied the asymptotics of the Cox partial likelihood; we mention, in particular, the result of Lin and Wei (1989) on asymptotic normality of the maximum partial likelihood estimator for the Cox model under misspecification.

The generalized posterior $\pi_n(\theta) \propto \mathcal{L}_n^{\mathrm{Cox}}(\theta)\pi(\theta)$ based on the Cox partial likelihood has been considered by several authors (e.g., Raftery et al., 1996; Sinha et al., 2003; Kim and Kim, 2009; Ventura and Racugno, 2016). Sinha et al. (2003) show that π_n approximates the standard posterior under a semiparametric Bayesian model, extending the results of Kalbfleisch (1978). Here, we provide sufficient conditions for π_n to exhibit concentration, asymptotic normality, and an asymptotically correct Laplace approximation.

Theorem 7.6. *Suppose* $(X, Y, Z), (X_1, Y_1, Z_1), (X_2, Y_2, Z_2), \ldots$ *are i.i.d., where* $X \in \mathcal{X} \subseteq \mathbb{R}^D$, $Y \geq 0$, *and* $Z \in \{0, 1\}$. *Define* $f(\theta) = \mathrm{E}\big(h_Y(\theta)Z\big) - \theta^{\mathsf{T}}\mathrm{E}(XZ)$ *for* $\theta \in \Theta := \mathbb{R}^D$ *where* $h_y(\theta) = \log \mathrm{E}\big(\exp(\theta^{\mathsf{T}}X)\mathbb{1}(Y \geq y)\big)$. *Assume:*

1. $\mathcal{X}$ *is bounded,*

2. *the c.d.f. of* Y *is continuous on* $\mathbb{R}$,

3. $\mathbb{P}(Z = 1) > 0$ *and* $\mathrm{Var}(a^{\mathsf{T}}X) > 0$ *for all nonzero* $a \in \mathbb{R}^D$,

4. $\mathbb{P}(Y \geq y \mid X = x) > 0$ *for all* $x \in \mathcal{X}$, $y \geq 0$, *and*

5. $f'(\theta_0) = 0$ *for some* $\theta_0 \in \mathbb{R}^D$.

Then for any open ball E *such that* $\theta_0 \in E$, $f_n(\theta) := -\frac{1}{n}\log \mathcal{L}_n^{\mathrm{Cox}}(\theta) - \frac{1}{n}\sum_{i=1}^{n} Z_i \log n$ *satisfies the conditions of Theorem 3.2 with probability 1.*

See Section S5 for the proof. Note that $\exp(-nf_n(\theta)) \propto \mathcal{L}_n^{\mathrm{Cox}}(\theta)$ since $\frac{1}{n}\sum_{i=1}^{n} Z_i \log n$ does not depend on θ; the purpose of introducing this term is so that f_n converges.

7.6 Median-based posterior for a location parameter

Suppose we wish to perform robust Bayesian inference for the parameter θ of a location family model $G_\theta(x) = G(x - \theta)$ where G is a cumulative distribution function (c.d.f.) on $\mathbb{R}$. If G is misspecified, then the posterior on θ can be poorly behaved, and may even fail to converge at all. For instance, if G_θ is the c.d.f. of $\mathcal{N}(\theta, \sigma^2)$ and the data are $X_1, X_2, \ldots$ i.i.d. $\sim \mathrm{Cauchy}(0, 1)$, then the posterior on θ is concentrated near $\frac{1}{n}\sum_{i=1}^{n} X_i$ when n is large, but $\frac{1}{n}\sum_{i=1}^{n} X_i \sim \mathrm{Cauchy}(0, 1)$; thus, the posterior does not converge to any fixed value.

Doksum and Lo (1990) propose to use the conditional distribution of θ given the sample median (or some other robust estimate of location) to perform robust Bayesian inference for θ. More precisely, let $M(x_{1:n})$ be a sample median of $x_{1:n} = (x_1, \ldots, x_n)$ and assume G_θ has a density g_θ. Then when n is odd,

$$
\begin{aligned}
p\big(\theta \mid M(X_{1:n}) = m\big) &\propto \pi(\theta)\, p(M(X_{1:n}) = m \mid \theta) \\
&\propto \pi(\theta)g_\theta(m)G_\theta(m)^{(n-1)/2}\big(1 - G_\theta(m)\big)^{(n-1)/2} \\
&= \pi(\theta)\exp\Big(\tfrac{1}{2}(n-1)\log G(m - \theta)(1 - G(m - \theta)) + \log g_\theta(m)\Big)
\end{aligned}
$$

where π is the prior on θ. Here, the conditional densities are under the model in which $\theta \sim \pi$ and $X_1, \ldots, X_n | \theta$ i.i.d. $\sim G_\theta$. Doksum and Lo (1990) show that $p(\theta \mid M(X_{1:n}) = M(x_{1:n}))$

20

and generalizations thereof have desirable properties as robust posteriors for θ; in particular, they provide consistency and asymptotic normality results.

With this as motivation, consider the generalized posterior $\pi_n(\theta) \propto \pi(\theta) \exp(-n f_n(\theta))$ where $f_n(\theta) = -\frac{1}{2} \log G(m_n - \theta)(1 - G(m_n - \theta))$ and $m_n = M(x_{1:n})$; this approximates $p(\theta \mid M(X_{1:n}) = m_n)$ and is somewhat simpler to analyze. The following theorem strengthens the Doksum and Lo (1990) asymptotic normality result by showing convergence in total variation distance, rather than convergence in the weak topology. Further, our conditions are simpler, but we do assume greater regularity of G and we only consider the median.

Theorem 7.7. *Suppose $G : \mathbb{R} \to (0, 1)$ is a c.d.f. such that G''' exists and is continuous, $G(-x) = 1 - G(x)$ for all $x \in \mathbb{R}$, $(\log G)''(x) \leq 0$ for all $x \in \mathbb{R}$, and $(\log G)''(0) < 0$. If $\theta_0 \in \mathbb{R}$ and $m_1, m_2, \ldots \in \mathbb{R}$ such that $\theta_0 = \lim_{n \to \infty} m_n$, then for any open ball E containing θ_0, $f_n(\theta) := -\frac{1}{2} \log G(m_n - \theta)(1 - G(m_n - \theta))$ satisfies the conditions of Theorem 3.2 on $\mathbb{R}$.*

Proof. By the chain rule, $f_n(\theta)$ has a continuous third derivative since $\log(x)$ and $G(x)$ have continuous third derivatives and $G(x) \in (0, 1)$. Define $f(\theta) = -\frac{1}{2} \log G(\theta_0 - \theta)(1 - G(\theta_0 - \theta))$ for $\theta \in \mathbb{R}$. Then for all $\theta \in \mathbb{R}$, $f_n(\theta) \to f(\theta)$ as $n \to \infty$ since $m_n \to \theta_0$, $\log(x)$ and $G(x)$ are continuous, and $G(x) \in (0, 1)$. Further,

$$f(\theta) = -\tfrac{1}{2} \log G(\theta_0 - \theta) - \tfrac{1}{2} \log G(\theta - \theta_0),$$
$$f'(\theta) = \tfrac{1}{2}(\log G)'(\theta_0 - \theta) - \tfrac{1}{2}(\log G)'(\theta - \theta_0),$$
$$f''(\theta) = -\tfrac{1}{2}(\log G)''(\theta_0 - \theta) - \tfrac{1}{2}(\log G)''(\theta - \theta_0).$$

Thus, $f'(\theta_0) = 0$ and $f''(\theta_0) = -(\log G)''(0) > 0$. Similarly, $f_n''(\theta) = -\frac{1}{2}(\log G)''(m_n - \theta) - \frac{1}{2}(\log G)''(\theta - m_n) \geq 0$ since $(\log G)''(x) \leq 0$. Thus, f_n is convex. Finally, for any bounded open interval E containing θ_0, (f_n''') is uniformly bounded on E by Proposition 7.8 with $h(\theta, s) = -\frac{1}{2} \log G(s - \theta) G(\theta - s)$, $K = \bar{E}$, and $S = [\inf m_n, \sup m_n] \subseteq \mathbb{R}$. $\square$

In cases where $f_n(\theta) = h(\theta, s_n)$ for some finite-dimensional statistic s_n, the following simple proposition can make it easy to verify the uniform boundedness condition.

Proposition 7.8. *Let $K \subseteq \mathbb{R}^D$ and $S \subseteq \mathbb{R}^d$ be compact sets. Suppose $f_n(\theta) = h(\theta, s_n)$ for $\theta \in K$, $n \in \mathbb{N}$, where $h : K \times S \to \mathbb{R}$ and $s_1, s_2, \ldots \in S$. If $(\partial^3 h / \partial \theta_i \partial \theta_j \partial \theta_k)(\theta, s)$ exists and is continuous on $K \times S$ for all $i, j, k \in \{1, \ldots, D\}$, then (f_n''') is uniformly bounded on K.*

Proof. Let $h'''(\theta, s)$ denote the tensor of third derivatives with respect to θ, and let $c = \sup\{\|h'''(\theta, s)\| : \theta \in K, s \in S\}$. For all $\theta \in K$, $n \in \mathbb{N}$, we have $\|f_n'''(\theta)\| = \|h'''(\theta, s_n)\| \leq c$, and $c < \infty$ since $(\theta, s) \mapsto \|h'''(\theta, s)\|$ is continuous and $K \times S$ is compact. $\square$

Acknowledgments

References

D. H. Ackley, G. E. Hinton, and T. J. Sejnowski. A learning algorithm for Boltzmann machines. *Cognitive Science*, 9(1):147–169, 1985.

S. Banerjee, B. P. Carlin, and A. E. Gelfand. *Hierarchical Modeling and Analysis for Spatial Data*. Chapman and Hall/CRC, 2014.

J. M. Bernardo and A. F. Smith. *Bayesian Theory*. John Wiley & Sons, 2000.

J. Besag. Statistical analysis of non-lattice data. *Journal of the Royal Statistical Society: Series D (The Statistician)*, 24(3):179–195, 1975.

P. J. Bickel and B. J. Kleijn. The semiparametric Bernstein–von Mises theorem. *The Annals of Statistics*, 40(1):206–237, 2012.

P. G. Bissiri, C. C. Holmes, and S. G. Walker. A general framework for updating belief distributions. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 78(5):1103–1130, 2016.

N. A. Bochkina and P. J. Green. The Bernstein–von Mises theorem and nonregular models. *The Annals of Statistics*, 42(5):1850–1878, 2014.

L. Bouranis, N. Friel, and F. Maire. Efficient Bayesian inference for exponential random graph models by correcting the pseudo-posterior distribution. *Social Networks*, 50:98–108, 2017.

I. Castillo and R. Nickl. Nonparametric Bernstein–von Mises theorems in Gaussian white noise. *The Annals of Statistics*, 41(4):1999–2028, 2013.

I. Castillo and J. Rousseau. A General Bernstein–von Mises Theorem in semiparametric models. *arXiv preprint arXiv:1305.4482*, 2013.

F. Comets. On consistency of a class of estimators for exponential families of Markov random fields on the lattice. *The Annals of Statistics*, pages 455–468, 1992.

D. R. Cox. Regression models and life-tables. *Journal of the Royal Statistical Society: Series B (Methodological)*, 34(2):187–202, 1972.

D. R. Cox. Partial likelihood. *Biometrika*, 62(2):269–276, 1975.

K. A. Doksum and A. Y. Lo. Consistent and robust Bayes procedures for location based on partial information. *The Annals of Statistics*, 18(1):443–453, 1990.

R. M. Dudley. *Real Analysis and Probability*. Cambridge University Press, 2002.

D. B. Dunson and J. A. Taylor. Approximate Bayesian inference for quantiles. *Nonparametric Statistics*, 17(3):385–400, 2005.

B. Efron. The efficiency of Cox's likelihood function for censored data. *Journal of the American Statistical Association*, 72(359):557–565, 1977.

G. B. Folland. *Real Analysis: Modern Techniques and Their Applications*. John Wiley & Sons, 2013.

N. Friel. Bayesian inference for Gibbs random fields using composite likelihoods. In *Proceedings of the Winter Simulation Conference*. Winter Simulation Conference, 2012.

S. Geman and D. Geman. Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 6(6):721, 1984.

S. Geman and C. Graffigne. Markov random field image models and their applications to computer vision. In *Proceedings of the International Congress of Mathematicians*, pages 1496–1517. Berkeley, CA, 1986.

J. K. Ghosh and R. Ramamoorthi. *Bayesian Nonparametrics*. Springer, 2003.

B. Gidas. Consistency of maximum likelihood and pseudo-likelihood estimators for Gibbs distributions. In *Stochastic Differential Systems, Stochastic Control Theory and Applications*, pages 129–145. Springer, 1988.

L. Greco, W. Racugno, and L. Ventura. Robust likelihood functions in Bayesian inference. *Journal of Statistical Planning and Inference*, 138(5):1258–1270, 2008.

G. E. Hinton and T. J. Sejnowski. Optimal perceptual inference. In *Proceedings of the IEEE conference on Computer Vision and Pattern Recognition*, pages 448–453, 1983.

N. L. Hjort and D. Pollard. Asymptotics for minimisers of convex processes. *Universitetet i Oslo. Matematisk Institutt*, 1993.

P. D. Hoff. Extending the rank likelihood for semiparametric copula estimation. *The Annals of Applied Statistics*, 1(1):265–283, 2007.

A. Hyvärinen. Consistency of pseudolikelihood estimation of fully visible Boltzmann machines. *Neural Computation*, 18(10):2283–2292, 2006.

P. E. Jacob, L. M. Murray, C. C. Holmes, and C. P. Robert. Better together? Statistical learning in models made of modules. *arXiv preprint arXiv:1708.08719*, 2017.

J. L. Jensen and H. R. Künsch. On asymptotic normality of pseudo likelihood estimates for pairwise interaction processes. *Annals of the Institute of Statistical Mathematics*, 46(3): 475–486, 1994.

W. Jiang and M. A. Tanner. Gibbs posterior for variable selection in high-dimensional classification and data mining. *The Annals of Statistics*, 36(5):2207–2231, 2008.

J. D. Kalbfleisch. Non-parametric Bayesian analysis of survival time data. *Journal of the Royal Statistical Society: Series B (Methodological)*, 40(2):214–221, 1978.

O. Kallenberg. *Foundations of Modern Probability*. Springer, 2002.

Y. Kim and D. Kim. Bayesian partial likelihood approach for tied observations. *Journal of Statistical Planning and Inference*, 139(2):469–477, 2009.

Y. Kim and J. Lee. A Bernstein–von Mises theorem in the nonparametric right-censoring model. *The Annals of Statistics*, 32(4):1492–1512, 2004.

B. Kleijn and A. Van der Vaart. The Bernstein-von-Mises theorem under misspecification. *Electronic Journal of Statistics*, 6:354–381, 2012.

R. Lang. A note on the measurability of convex sets. *Archiv der Mathematik*, 47(1):90–92, 1986.

M. Lavine. On an approximate likelihood for quantiles. *Biometrika*, 82(1):220–222, 1995.

N. A. Lazar. Bayesian empirical likelihood. *Biometrika*, 90(2):319–326, 2003.

H. Leahu. On the Bernstein-von Mises phenomenon in the Gaussian white noise model. *Electronic Journal of Statistics*, 5:373–404, 2011.

J. R. Lewis, S. N. MacEachern, and Y. Lee. Bayesian restricted likelihood methods. *Technical report 878, The Ohio State University*, 2014.

G. Liang and B. Yu. Maximum pseudo likelihood estimation in network tomography. *IEEE Transactions on Signal Processing*, 51(8):2043–2053, 2003.

D. Y. Lin and L.-J. Wei. The robust inference for the Cox proportional hazards model. *Journal of the American Statistical Association*, 84(408):1074–1078, 1989.

B. G. Lindsay. Composite likelihood methods. *Contemporary Mathematics*, 80(1):221–239, 1988.

F. Liu, M. Bayarri, and J. Berger. Modularization in Bayesian analysis, with emphasis on analysis of computer models. *Bayesian Analysis*, 4(1):119–150, 2009.

S. Mase. Consistency of the maximum pseudo-likelihood estimator of continuous state space Gibbsian processes. *The Annals of Applied Probability*, 5(3):603–612, 1995.

J. W. Miller and M. T. Harrison. Inconsistency of Pitman–Yor process mixtures for the number of components. *The Journal of Machine Learning Research*, 15(1):3333–3370, 2014.

G. Molenbergs and G. Verbeke. *Models for Discrete Longitudinal Data*. Springer Science & Business Media, 2005.

M. Panov and V. Spokoiny. Finite sample Bernstein–von Mises theorem for semiparametric problems. *Bayesian Analysis*, 10(3):665–710, 2015.

F. Pauli, W. Racugno, and L. Ventura. Bayesian composite marginal likelihoods. *Statistica Sinica*, pages 149–164, 2011.

A. Pettitt. Likelihood based inference using signed ranks for matched pairs. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 45(2):287–296, 1983.

A. E. Raftery, D. Madigan, and C. T. Volinsky. Accounting for model uncertainty in survival analysis improves predictive performance. *Bayesian Statistics*, 5:323–349, 1996.

M. Ribatet, D. Cooley, and A. C. Davison. Bayesian inference from composite likelihoods, with an application to spatial extremes. *Statistica Sinica*, pages 813–845, 2012.

R. T. Rockafellar. *Convex Analysis*. Princeton University Press, 1970.

W. Rudin. *Principles of Mathematical Analysis*. McGraw-Hill New York, 1976.

T. Rydén and D. Titterington. Computational Bayesian analysis of hidden Markov models. *Journal of Computational and Graphical Statistics*, 7(2):194–211, 1998.

L. Schwartz. On Bayes procedures. *Zeitschrift für Wahrscheinlichkeitstheorie und verwandte Gebiete*, 4(1):10–26, 1965.

X. Shen. Asymptotic normality of semiparametric and nonparametric posterior distributions. *Journal of the American Statistical Association*, 97(457):222–235, 2002.

D. Sinha, J. G. Ibrahim, and M.-H. Chen. A Bayesian justification of Cox's partial likelihood. *Biometrika*, 90(3):629–641, 2003.

E. L. Smith and A. G. Stephenson. An extended Gaussian max-stable process model for spatial extremes. *Journal of Statistical Planning and Inference*, 139(4):1266–1275, 2009.

J. Stoehr and N. Friel. Calibration of conditional composite likelihood for Bayesian inference on Gibbs random fields. In *Proceedings of the 18th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 38. JMLR, 2015.

A. W. Van der Vaart. *Asymptotic Statistics*. Cambridge University Press, 2000.

C. Varin, N. Reid, and D. Firth. An overview of composite likelihood methods. *Statistica Sinica*, pages 5–42, 2011.

L. Ventura and W. Racugno. Pseudo-likelihoods for Bayesian inference. In *Topics on Methodological and Applied Statistical Inference*, pages 205–220. Springer, 2016.

L. Ventura, S. Cabras, and W. Racugno. Default prior distributions from quasi-and quasi-profile likelihoods. *Journal of Statistical Planning and Inference*, 140(11):2937–2942, 2010.

M. R. Williams and T. D. Savitsky. Bayesian uncertainty estimation under complex sampling. *arXiv preprint arXiv:1807.11796*, 2018.

G. Yin. Bayesian generalized method of moments. *Bayesian Analysis*, 4(2):191–207, 2009.

T. Zhang. Information-theoretic upper and lower bounds for statistical estimation. *IEEE Transactions on Information Theory*, 52(4):1307–1321, 2006.

X. Zhou and S. C. Schmidler. Bayesian parameter estimation in Ising and Potts models: A comparative study with applications to protein modeling. *Technical report, Duke University*, 2009.

Supplementary material for "Asymptotic normality, concentration, and coverage of generalized posteriors"

S1 Proofs of concentration results

Proof of Theorem 2.2. Let $\varepsilon > 0$. Define $\mu_n(E) = \int_E e^{-nf_n(\theta)}\Pi(d\theta)$ for $E \subseteq \Theta$. Recall that $\mu_n(\Theta) = z_n < \infty$ by assumption. For any $\beta \in \mathbb{R}$,

$$1 - \Pi_n(A_\varepsilon) = \Pi_n(A_\varepsilon^c) = \frac{\mu_n(A_\varepsilon^c)}{\mu_n(\Theta)} = \frac{e^{n(f(\theta_0)+\beta)}\mu_n(A_\varepsilon^c)}{e^{n(f(\theta_0)+\beta)}\mu_n(\Theta)},$$

so prove the result, it suffices to show that for some β, the numerator is bounded and the denominator goes to ∞.

First, consider the numerator. Condition 3 implies that there exists $\beta > 0$ such that for all n sufficiently large, $\inf_{\theta \in A_\varepsilon^c} f_n(\theta) \geq f(\theta_0) + \beta$. Then for all n sufficiently large, for all $\theta \in A_\varepsilon^c$, we have $\exp\big(-n(f_n(\theta) - f(\theta_0) - \beta)\big) \leq 1$. Hence, for all n sufficiently large,

$$e^{n(f(\theta_0)+\beta)}\mu_n(A_\varepsilon^c) = \int_{A_\varepsilon^c} \exp\big(-n(f_n(\theta) - f(\theta_0) - \beta)\big)\Pi(d\theta) \leq \int_{A_\varepsilon^c} \Pi(d\theta) \leq 1.$$

Now, consider the denominator. For any $\theta \in A_{\beta/2}$, $f_n(\theta) - f(\theta_0) - \beta \longrightarrow f(\theta) - f(\theta_0) - \beta < -\beta/2 < 0$, and thus, $\exp\big(-n(f_n(\theta) - f(\theta_0) - \beta)\big) \longrightarrow \infty$ as $n \to \infty$. Therefore, by Fatou's lemma,

$$\liminf_{n\to\infty} e^{n(f(\theta_0)+\beta)}\mu_n(A_{\beta/2}) = \liminf_{n\to\infty} \int_{A_{\beta/2}} \exp\big(-n(f_n(\theta) - f(\theta_0) - \beta)\big)\Pi(d\theta) = \infty$$

since $\Pi(A_{\beta/2}) > 0$. Hence, $e^{n(f(\theta_0)+\beta)}\mu_n(\Theta) \to \infty$ since $\mu_n(\Theta) \geq \mu_n(A_{\beta/2})$. $\square$

Lemma S1.1. *Suppose $\Theta \subseteq \mathbb{R}^D$, $E \subseteq \Theta$ is convex and open in $\mathbb{R}^D$, and $\theta_0 \in E$. Let $f_n : \Theta \to \mathbb{R}$ be convex, and assume $f_n \to f$ pointwise on E for some $f : E \to \mathbb{R}$.*

1. *If f' exists in a neighborhood of θ_0, $f'(\theta_0) = 0$, and $f''(\theta_0)$ exists and is positive definite, then $f(\theta) > f(\theta_0)$ for all $\theta \in E \setminus \{\theta_0\}$.*

2. *If $f(\theta) > f(\theta_0)$ for all $\theta \in E \setminus \{\theta_0\}$, then $\liminf_n \inf_{\theta \in \Theta \setminus B_\varepsilon(\theta_0)} f_n(\theta) > f(\theta_0)$ for any $\varepsilon > 0$.*

Proof. (1) As the pointwise limit of convex functions on a convex open set, f is convex on E (Rockafellar, 1970, 10.8). Let $R > 0$ such that $f'(\theta)$ exists for all $\theta \in B_R(\theta_0)$. Let $u \in \mathbb{R}^D$ with $|u| = 1$, and define $g(r) = f(\theta_0 + ru)$ for $r \in [0, R)$. Then $g'(r) = f'(\theta_0 + ru)^\mathsf{T} u$ and $g''(0) = u^\mathsf{T} f''(\theta_0)u$. Since

$$\frac{g'(r)}{r} = \frac{g'(r) - g'(0)}{r} \xrightarrow[r\to 0]{} g''(0) = u^\mathsf{T} f''(\theta_0)u > 0,$$

then $g'(r) > 0$ for all $r > 0$ sufficiently small, say, all $r \in (0, \varepsilon]$. Then for any $s \in (0, \varepsilon]$, we have

$$f(\theta_0 + su) - f(\theta_0) = g(s) - g(0) = \int_0^s g'(r)dr > 0. \tag{S1}$$

Meanwhile, for any $s > \varepsilon$ such that $\theta_0 + su \in E$, we have

$$\frac{1}{s}(f(\theta_0 + su) - f(\theta_0)) \geq \frac{1}{\varepsilon}(f(\theta_0 + \varepsilon u) - f(\theta_0)) > 0$$

by the convexity of f and by Equation S1 with $s = \varepsilon$. Hence, for any $s > 0$ such that $\theta_0 + su \in E$, $f(\theta_0 + su) > f(\theta_0)$. Since u is arbitrary, the result follows.

(2) By Rockafellar (1970), 10.8, $f_n \to f$ uniformly on any compact subset of E, and f is convex on E. Further, f is continuous on E, as a convex function on a convex open set (Rockafellar, 1970, Theorem 10.1). Let $\varepsilon > 0$ small enough that the ε-sphere $S_\varepsilon = \{\theta \in \mathbb{R}^D : |\theta - \theta_0| = \varepsilon\}$ is contained in E. Let $\alpha_n = \inf_{\theta \in S_\varepsilon} f_n(\theta) - f_n(\theta_0)$ and $\alpha = \inf_{\theta \in S_\varepsilon} f(\theta) - f(\theta_0)$. By uniform convergence, $\alpha_n \to \alpha$. Note that $\alpha > 0$, as the minimum of the continuous positive function $f(\theta) - f(\theta_0)$ on the compact set S_ε. For any $\theta \in \Theta \setminus B_\varepsilon(\theta_0)$, letting ξ_θ be the point of S_ε on the line from θ to θ_0, we have, by the convexity of f_n,

$$f_n(\theta) - f_n(\theta_0) \geq |\theta - \theta_0| \frac{f_n(\xi_\theta) - f_n(\theta_0)}{|\xi_\theta - \theta_0|} \geq \alpha_n$$

whenever $\alpha_n \geq 0$, since $|\theta - \theta_0| \geq |\xi_\theta - \theta_0|$. Since $\alpha_n \to \alpha > 0$, then for all n sufficiently large, for all $\theta \in \Theta \setminus B_\varepsilon(\theta_0)$, $f_n(\theta) \geq f_n(\theta_0) + \alpha_n \longrightarrow f(\theta_0) + \alpha$. Therefore, $\liminf_n \inf_{\theta \in \Theta \setminus B_\varepsilon(\theta_0)} f_n(\theta) \geq f(\theta_0) + \alpha > f(\theta_0)$. Note that this also implies the same inequality for any $\varepsilon' > \varepsilon$. $\square$

Proof of Theorem 2.3.

   (**Part 1**) Defining A_ε as in Theorem 2.2, it suffices to show that

   (a) for any $\varepsilon > 0$ there exists $\delta > 0$ such that $A_\delta \subseteq N_\varepsilon$, and

   (b) for any $\delta > 0$ there exists $\varepsilon' > 0$ such that $N_{\varepsilon'} \subseteq A_\delta$,

since for any $\varepsilon > 0$, choosing δ by (a), we have $\Pi_n(N_\varepsilon) \geq \Pi_n(A_\delta)$; meanwhile, for any $\delta > 0$, choosing ε' by (b), we have $\Pi(A_\delta) \geq \Pi(N_{\varepsilon'}) > 0$ and $\liminf_n \inf_{\theta \in A_\delta^c} f_n(\theta) \geq \liminf_n \inf_{\theta \in N_{\varepsilon'}^c} f_n(\theta) > f(\theta_0)$, and hence, by Theorem 2.2, $\Pi_n(A_\delta) \to 1$.

   (a) Let $\varepsilon > 0$. Pointwise convergence and the liminf condition imply $\inf_{\theta \in N_\varepsilon^c} f(\theta) > f(\theta_0)$, hence, letting $\delta = \inf_{\theta \in N_\varepsilon^c} f(\theta) - f(\theta_0)$, we have $\delta > 0$ and $A_\delta \subseteq N_\varepsilon$.

   (b) Let $\delta > 0$. By the continuity of f at θ_0, choose $\varepsilon' > 0$ such that $|f(\theta) - f(\theta_0)| < \delta$ for all $\theta \in N_{\varepsilon'}$. Then for any $\theta \in N_{\varepsilon'}$, $f(\theta) < f(\theta_0) + \delta$, hence, $\theta \in A_\delta$.

   (**Part 2**) We show that 2 implies 1. By Lemma S6.1, $f_n \to f$ uniformly on K. Consequently, $f|_K$ is continuous, as the uniform limit of continuous functions (Rudin, 1976, 7.12). In particular, f is continuous at θ_0, since θ_0 is an interior point of K. For any $\varepsilon > 0$,

$$\liminf_n \inf_{\theta \in K \setminus N_\varepsilon} f_n(\theta) = \inf_{\theta \in K \setminus N_\varepsilon} f(\theta) > f(\theta_0),$$

the first step holding since $f_n \to f$ uniformly on K, and the second step since $f|_K$ is continuous, $K \setminus N_\varepsilon$ is compact, and $f(\theta) > f(\theta_0)$ for all $\theta \in K \setminus \{\theta_0\}$. Therefore, since $N_\varepsilon^c \subseteq (K \setminus N_\varepsilon) \cup K^c$,

$$\liminf_n \inf_{\theta \in N_\varepsilon^c} f_n(\theta) \geq \liminf_n \min \left\{ \inf_{\theta \in K \setminus N_\varepsilon} f_n(\theta), \inf_{\theta \in K^c} f_n(\theta) \right\} > f(\theta_0).$$

(Part 3) We show that 3 implies 1. Denote $B_\varepsilon = \{\theta \in \mathbb{R}^D : |\theta - \theta_0| < \varepsilon\}$. Let $r > 0$ small enough that $B_r \subseteq \Theta$. As the pointwise limit of convex functions, f is convex, and thus, it is continuous on B_r (Rockafellar, 1970, 10.1). By Lemma S1.1 with $E = B_r$, in either case (a) or (b), we have

$$\liminf_n \inf_{\theta \in \Theta \setminus B_\varepsilon} f_n(\theta) > f(\theta_0)$$

for any $\varepsilon > 0$. Since $\Theta \setminus B_\varepsilon = \Theta \setminus N_\varepsilon = N_\varepsilon^c$, this proves the result. $\qquad\square$

S2 Proofs of asymptotic normality results

Lemma S2.1. *Let $\theta_n \in \mathbb{R}^D$ such that $\theta_n \to \theta_0$ for some $\theta_0 \in \mathbb{R}^D$, let π_n be a density with respect to Lebesgue measure on $\mathbb{R}^D$, and let q_n be the density of $\sqrt{n}(\theta - \theta_n)$ when $\theta \sim \pi_n$. If $\int |q_n(x) - q(x)| dx \longrightarrow 0$ for some probability density q, then π_n concentrates at θ_0.*

Proof. Let Π_n, Q_n, and Q denote the probability measures corresponding to π_n, q_n, and q, respectively. For any $\varepsilon > 0$ and $\delta > 0$,

$$Q_n(B_\delta(0)) = \Pi_n(B_{\delta/\sqrt{n}}(\theta_n)) \leq \Pi_n(B_\varepsilon(\theta_0))$$

for all n sufficiently large. Hence, since $Q_n \to Q$ in total variation,

$$Q(B_\delta(0)) = \lim_n Q_n(B_\delta(0)) \leq \liminf_n \Pi_n(B_\varepsilon(\theta_0)).$$

Taking the limit as $\delta \to \infty$ shows that $\lim_n \Pi_n(B_\varepsilon(\theta_0)) = 1$. $\qquad\square$

Proof of Theorem 3.1. Note that $q_n(x) = \pi_n(\theta_n + x/\sqrt{n}) n^{-D/2}$. Define

$$\begin{aligned}
g_n(x) &= \exp\big(-n[f_n(\theta_n + x/\sqrt{n}) - f_n(\theta_n)]\big)\pi(\theta_n + x/\sqrt{n}) \\
&= q_n(x) e^{n f_n(\theta_n)} n^{D/2} z_n,
\end{aligned} \tag{S2}$$

recalling that $z_n < \infty$ by assumption, and define

$$g_0(x) = \exp(-\tfrac{1}{2} x^{\mathsf{T}} H_0 x) \pi(\theta_0).$$

Let $\alpha \in (0, \lambda)$, where λ is the smallest eigenvalue of H_0. Let $\varepsilon > 0$ small enough that $\varepsilon < \alpha/(2c_0)$, $\varepsilon < \varepsilon_0$, and $\pi(\theta) \leq 2\pi(\theta_0)$ for all $\theta \in B_{2\varepsilon}(\theta_0)$ (which we can do since π is continuous at θ_0). Let $\delta = \liminf_n \inf_{\theta \in B_\varepsilon(\theta_n)^c} \big(f_n(\theta) - f_n(\theta_n)\big)$, noting that $\delta > 0$ by assumption. Letting $A_n = H_n - \alpha I$ and $A_0 = H_0 - \alpha I$, define

$$h_n(x) = \begin{cases} \exp(-\tfrac{1}{2} x^{\mathsf{T}} A_n x) 2\pi(\theta_0) & \text{if } |x| < \varepsilon\sqrt{n}, \\ e^{-n\delta/2} \pi(\theta_n + x/\sqrt{n}) & \text{if } |x| \geq \varepsilon\sqrt{n}, \end{cases}$$

$$h_0(x) = \exp(-\tfrac{1}{2} x^{\mathsf{T}} A_0 x) 2\pi(\theta_0).$$

We will show that

(a) $g_n \to g_0$ and $h_n \to h_0$ pointwise,

(b) $\int h_n \to \int h_0$,

(c) $g_n = |g_n| \le h_n$ for all n sufficiently large, and

(d) $g_n, g_0, h_n, h_0 \in L^1$ for all n sufficiently large.

By the generalized dominated convergence theorem, this will imply that $\int g_n \to \int g_0$ and $\int |g_n - g_0| \to 0$ (e.g., Folland, 2013, exercises 2.20, 2.21). Supposing this for the moment, we show how the result follows. Since $\int q_n = 1$, by Equation S2 we have

$$e^{nf_n(\theta_n)} n^{D/2} z_n = \int g_n \longrightarrow \int g_0 = \pi(\theta_0) \frac{(2\pi)^{D/2}}{|H_0|^{1/2}}, \tag{S3}$$

where $|H_0| = |\det H_0|$, and hence,

$$z_n \sim \frac{e^{-nf_n(\theta_n)}\pi(\theta_0)}{|H_0|^{1/2}} \left(\frac{2\pi}{n}\right)^{D/2}$$

as $n \to \infty$; this proves Equation 3. For any $a_n \to a \in \mathbb{R}$, we have $\int |a_n g_n - a g_0| \to 0$ since

$$\int |a_n g_n - a g_0| \le \int |a_n g_n - a_n g_0| + \int |a_n g_0 - a g_0| = |a_n| \int |g_n - g_0| + |a_n - a| \int |g_0| \longrightarrow 0.$$

Thus, letting $1/a_n = e^{nf_n(\theta_n)} n^{D/2} z_n$ and $1/a = \pi(\theta_0)\frac{(2\pi)^{D/2}}{|H_0|^{1/2}}$, we have $a_n \to a$ by Equation S3, and thus,

$$\int \left| q_n(x) - \frac{|H_0|^{1/2}}{(2\pi)^{D/2}} \exp(-\tfrac{1}{2}x^\mathsf{T} H_0 x) \right| dx \longrightarrow 0,$$

proving Equation 4. Equation 2 (concentration at θ_0) follows by Lemma S2.1, since $\theta_n \to \theta_0$. It remains to show (a)–(d) above.

(a) Fix $x \in \mathbb{R}^D$. First, consider h_n. For all n sufficiently large, $|x| < \varepsilon\sqrt{n}$, and thus,

$$h_n(x) = \exp(-\tfrac{1}{2}x^\mathsf{T} A_n x) 2\pi(\theta_0) \longrightarrow \exp(-\tfrac{1}{2}x^\mathsf{T} A_0 x) 2\pi(\theta_0) = h_0(x)$$

since $A_n \to A_0$. Now, for g_n, first note that $\pi(\theta_n + x/\sqrt{n}) \to \pi(\theta_0)$ since π is continuous at θ_0 and $\theta_n \to \theta_0$, $x/\sqrt{n} \to 0$. By the assumed representation of f_n (Equation 1),

$$n(f_n(\theta_n + x/\sqrt{n}) - f_n(\theta_n)) = \tfrac{1}{2}x^\mathsf{T} H_n x + n r_n(x/\sqrt{n}) \longrightarrow \tfrac{1}{2}x^\mathsf{T} H_0 x$$

since $H_n \to H_0$ and for all n sufficiently large (to ensure that $|x/\sqrt{n}| < \varepsilon_0$ and the assumed bound on r_n holds),

$$|n r_n(x/\sqrt{n})| \le n c_0 |x/\sqrt{n}|^3 = c_0 |x|^3 / \sqrt{n} \to 0 \tag{S4}$$

as $n \to \infty$. Hence, $g_n(x) \to g_0(x)$.

(b) By the definition of h_n, letting $B_n = B_{\varepsilon\sqrt{n}}(0)$,

$$\int h_n = \int_{B_n} \exp(-\tfrac{1}{2}x^\mathsf{T} A_n x) 2\pi(\theta_0) dx + \int_{B_n^c} e^{-n\delta/2} \pi(\theta_n + x/\sqrt{n}) dx.$$

Since $A_n \to A_0$ and A_0 is positive definite, then for all n sufficiently large, A_n is also positive definite and the first term equals

$$2\pi(\theta_0)\frac{(2\pi)^{D/2}}{|A_n|^{1/2}}\mathbb{P}(|A_n^{-1/2}Z| < \varepsilon\sqrt{n}) \longrightarrow 2\pi(\theta_0)\frac{(2\pi)^{D/2}}{|A_0|^{1/2}} = \int h_0$$

where $Z \sim \mathcal{N}(0, I)$. The second term goes to zero, since it is nonnegative and upper bounded by

$$\int_{\mathbb{R}^D} e^{-n\delta/2}\pi(\theta_n + x/\sqrt{n})dx = e^{-n\delta/2}n^{D/2} \longrightarrow 0,$$

using the fact that $\pi(\theta_n + x/\sqrt{n})n^{-D/2}$ is the density of $X = \sqrt{n}(\theta - \theta_n)$ when $\theta \sim \pi$.

(c) For all n sufficiently large, $|\theta_n - \theta_0| < \varepsilon$, the bound on r_n applies, and $\inf_{\theta \in B_\varepsilon(\theta_n)^c} f_n(\theta) - f_n(\theta_n) > \delta/2$. Let n large enough that these hold, and let $x \in \mathbb{R}^D$. If $|x| \geq \varepsilon\sqrt{n}$, then $f_n(\theta_n + x/\sqrt{n}) - f_n(\theta_n) > \delta/2$, and thus,

$$g_n(x) \leq e^{-n\delta/2}\pi(\theta_n + x/\sqrt{n}) = h_n(x).$$

Meanwhile, if $|x| < \varepsilon\sqrt{n}$, then $\pi(\theta_n + x/\sqrt{n}) \leq 2\pi(\theta_0)$ (by our choice of ε, since $|(\theta_n + x/\sqrt{n}) - \theta_0| \leq |\theta_n - \theta_0| + |x/\sqrt{n}| < 2\varepsilon$), and

$$n(f_n(\theta_n + x/\sqrt{n}) - f_n(\theta_n)) = \tfrac{1}{2}x^\mathsf{T}H_n x + nr_n(x/\sqrt{n}) \geq \tfrac{1}{2}x^\mathsf{T}H_n x - \tfrac{1}{2}\alpha x^\mathsf{T}x = \tfrac{1}{2}x^\mathsf{T}A_n x$$

since $|nr_n(x/\sqrt{n})| \leq c_0|x|^3/\sqrt{n} \leq c_0\varepsilon|x|^2 \leq \tfrac{1}{2}\alpha|x|^2$, by the fact that $|x/\sqrt{n}| < \varepsilon < \varepsilon_0$ and $\varepsilon < \alpha/(2c_0)$. Therefore,

$$g_n(x) \leq \exp(-\tfrac{1}{2}x^\mathsf{T}A_n x)2\pi(\theta_0) = h_n(x).$$

(d) Since H_0 and A_0 are positive definite, $\int g_0$ and $\int h_0$ are finite. By (b) and (c), since $\int h_n \to \int h_0 < \infty$, we have $\int g_n \leq \int h_n < \infty$ for all n sufficiently large. Measurability of g_n and h_n follows from measurability of f_n and π. □

Proof of Theorem 3.2. First, we show that under case 2, the conditions for case 1 hold. By Lemma S1.1(1), $f(\theta) > f(\theta_0)$ for all $\theta \in E \setminus \{\theta_0\}$ since f' exists on E by Theorem 3.4. Letting $K = \overline{B_\varepsilon(\theta_0)}$ where $\varepsilon > 0$ is small enough that $K \subseteq E$, we have $\liminf_n \inf_{\theta \in \Theta \setminus K} f_n(\theta) > f(\theta_0)$ by Lemma S1.1(2). Thus, it suffices to prove the result under case 1.

Consider case 1. Extend π, f_n, and f to all of $\mathbb{R}^D$ by defining $\pi(\theta) = 0$ and $f(\theta) = f_n(\theta) = f(\theta_0) + 1$ for all $\theta \in \mathbb{R}^D \setminus \Theta$. Then all the conditions of Theorem 3.2 (under case 1) still hold with $\mathbb{R}^D$ in place of Θ. We will show that:

(a) (f_n) is equicontinuous on E, and $f_n''(\theta_0) \to f''(\theta_0)$ as $n \to \infty$,

(b) there exist $\theta_n \in E$ such that $\theta_n \to \theta_0$ and $f_n'(\theta_n) = 0$ for all n sufficiently large, and

(c) $f_n(\theta_n) \to f(\theta_0)$.

Assuming (a)–(c) for the moment, we show how the result follows. Letting $H_0 = f''(\theta_0)$, the conditions of Theorem 3.3 are satisfied, and thus, condition 1 of Theorem 3.1 is satisfied for all n sufficiently large. Condition 2 of Theorem 3.1 holds, since for all $\varepsilon > 0$,

$$\liminf_n \inf_{\theta \in B_\varepsilon(\theta_n)^c} (f_n(\theta) - f_n(\theta_n)) = \left(\liminf_n \inf_{\theta \in B_\varepsilon(\theta_n)^c} f_n(\theta) \right) - f(\theta_0)$$

$$\geq \left(\liminf_n \inf_{\theta \in B_{\varepsilon/2}(\theta_0)^c} f_n(\theta) \right) - f(\theta_0) > 0$$

the first step holding by (c), the second step since $\theta_n \to \theta_0$ and thus $B_{\varepsilon/2}(\theta_0) \subseteq B_\varepsilon(\theta_n)$ for all n sufficiently large, and the third step by the implication $2 \Rightarrow 1$ in Theorem 2.3. Thus, the conditions of Theorem 3.1 are satisfied (except possibly for some initial sequence of n's, which can be ignored since the conclusions are asymptotic in nature), establishing Equation 2 (concentration at θ_0), Equation 3 (the Laplace approximation), and Equation 4 (asymptotic normality). To complete the proof, we establish (a), (b), and (c).

(a) By Theorem 3.4, (f_n) is equi-Lipschitz (hence, equicontinuous) on E and $f_n'' \to f''$ uniformly on E.

(b) Let $\varepsilon > 0$ small enough that $S_\varepsilon \subseteq K$ where $S_\varepsilon = \{\theta \in \mathbb{R}^D : |\theta - \theta_0| = \varepsilon\}$. By Theorem 3.4, f is continuous on E (since f' exists on E). Thus, f attains its minimum on the compact set S_ε, and since $f(\theta) > f(\theta_0)$ on S_ε, we have $\inf_{\theta \in S_\varepsilon} f(\theta) > f(\theta_0)$. For each n, since f_n is continuous on E, its minimum over the set $\overline{B_\varepsilon(\theta_0)}$ is attained at one or more points; define θ_n^ε to be such a minimizer. Since $f_n \to f$ uniformly on E (by Theorem 3.4), then for all n sufficiently large, any such minimizer cannot be in S_ε (since $\inf_{\theta \in S_\varepsilon} f(\theta) > f(\theta_0)$). Hence, for all sufficiently small $\varepsilon > 0$, for all n sufficiently large, we have $\theta_n^\varepsilon \in B_\varepsilon(\theta_0)$ and (by Lemma S6.3) $f_n'(\theta_n^\varepsilon) = 0$.

Thus, we can choose a sequence $\varepsilon_n > 0$ such that (a) $\varepsilon_n \to 0$ and (b) for all n sufficiently large, $\theta_n^{\varepsilon_n} \in B_{\varepsilon_n}(\theta_0)$ and $f_n'(\theta_n^{\varepsilon_n}) = 0$. Therefore, letting $\theta_n = \theta_n^{\varepsilon_n}$, we have $\theta_n \to \theta_0$ and $f_n'(\theta_n) = 0$ for all n sufficiently large.

(c) We have $|f_n(\theta_n) - f(\theta_0)| \leq |f_n(\theta_n) - f_n(\theta_0)| + |f_n(\theta_0) - f(\theta_0)| \to 0$, the first term going to zero since $\theta_n \to \theta_0$ and (f_n) is equi-Lipschitz on E, and the second term since $f_n \to f$ pointwise. $\qquad\square$

For tensors $S, T \in \mathbb{R}^{D^3}$, define the inner product $\langle S, T \rangle = \sum_{i,j,k} S_{ijk} T_{ijk}$ (noting that this is just the dot product of the vectorized versions of S and T). For $x \in \mathbb{R}^D$, define $x^{\otimes 3} = x \otimes x \otimes x = \left(x_i x_j x_k \right)_{i,j,k=1}^D \in \mathbb{R}^{D^3}$, and note that $\|x^{\otimes 3}\| = |x|^3$.

Proof of Theorem 3.3. By Lemma S6.2, (f_n'') is equi-Lipschitz. Thus,

$$\|f_n''(\theta_n) - H_0\| \leq \|f_n''(\theta_n) - f_n''(\theta_0)\| + \|f_n''(\theta_0) - H_0\| \leq C|\theta_n - \theta_0| + \|f_n''(\theta_0) - H_0\| \longrightarrow 0,$$

and hence, $H_n \to H_0$. Let $C_0 = \sup_n \sup_{\theta \in E} \|f_n'''(\theta)\|$. Let n large enough that $f_n'(\theta_n) = 0$. For $\theta \in E$, by Taylor's theorem,

$$f_n(\theta) = f_n(\theta_n) + \tfrac{1}{2}(\theta - \theta_n)^{\mathsf{T}} f_n''(\theta_n)(\theta - \theta_n) + r_n(\theta - \theta_n)$$

where $r_n(\theta - \theta_n) = \tfrac{1}{6}\langle f_n'''(t_n(\theta)), (\theta - \theta_n)^{\otimes 3}\rangle$, and $t_n(\theta)$ is a point on the line between θ and θ_n. Then by Cauchy–Schwarz,

$$|r_n(\theta - \theta_n)| \leq \tfrac{1}{6}\|f_n'''(t_n(\theta))\|\|(\theta - \theta_n)^{\otimes 3}\| \leq \tfrac{1}{6}C_0\|(\theta - \theta_n)^{\otimes 3}\| = \tfrac{1}{6}C_0|\theta - \theta_n|^3. \qquad (S5)$$

Choose $\varepsilon_0 > 0$ small enough that $B_{2\varepsilon_0}(\theta_0) \subseteq E$, and choose $c_0 = C_0/6$. For all n sufficiently large, $|\theta_n - \theta_0| \leq \varepsilon_0$ and hence for all $x \in B_{\varepsilon_0}(0)$, we have $\theta_n + x \in B_{2\varepsilon_0}(\theta_0) \subseteq E$; thus, setting $\theta = \theta_n + x$ in Equation S5 yields $|r_n(x)| \leq c_0 |x|^3$. $\qquad \square$

S3 Proof of regular convergence theorem

Lemma S3.1. *Let $E \subseteq \mathbb{R}^D$ be open. If $f_n : E \to \mathbb{R}$ has continuous second derivatives, (f_n) is pointwise bounded, and (f_n'') is uniformly bounded, then (f_n') is pointwise bounded.*

Proof. Let $C = \sup\{\|f_n''(x)\| : n \in \mathbb{N}, x \in E\} < \infty$. Fix $x \in E$, and let $\varepsilon > 0$ small enough that $B_{2\varepsilon}(x) \subseteq E$. By Taylor's theorem, for any $u \in \mathbb{R}^D$ with $|u| = 1$,

$$f_n(x + \varepsilon u) = f_n(x) + \varepsilon f_n'(x)^\mathsf{T} u + \tfrac{1}{2}\varepsilon^2 u^\mathsf{T} f_n''(z) u$$

for some z on the line between x and $x + \varepsilon u$, and therefore,

$$|f_n'(x)^\mathsf{T} u| \leq (1/\varepsilon)|f_n(x + \varepsilon u) - f_n(x)| + \tfrac{1}{2}\varepsilon C$$

since $|u^\mathsf{T} f_n''(z) u| \leq \|f_n''(z)\| |u|^2 \leq C$. Thus, $\{f_n'(x)^\mathsf{T} u : n \in \mathbb{N}\}$ is bounded, for any u with $|u| = 1$. Applying this to each element of the standard basis, we have that $f_n'(x)$ is bounded. $\qquad \square$

Lemma S3.2. *Let $E \subseteq \mathbb{R}^D$ be open. If $f_n : E \to \mathbb{R}$ has continuous third derivatives, (f_n) is pointwise bounded, and (f_n''') is uniformly bounded, then (f_n'') is pointwise bounded.*

Proof. Let $C = \sup_n \sup_{x \in E} \|f_n'''(x)\| < \infty$. Fix $x \in E$, and let $\varepsilon > 0$ small enough that $\overline{B_\varepsilon(x)} \subseteq E$. By Taylor's theorem, for any $u \in \mathbb{R}^D$ with $|u| = 1$,

$$f_n(x + \varepsilon u) = f_n(x) + \varepsilon f_n'(x)^\mathsf{T} u + \tfrac{1}{2}\varepsilon^2 u^\mathsf{T} f_n''(x) u + \tfrac{1}{6}\varepsilon^3 \langle f_n'''(z^+), u^{\otimes 3}\rangle$$

for some z^+ on the line between x and $x + \varepsilon u$. Likewise,

$$f_n(x - \varepsilon u) = f_n(x) - \varepsilon f_n'(x)^\mathsf{T} u + \tfrac{1}{2}\varepsilon^2 u^\mathsf{T} f_n''(x) u - \tfrac{1}{6}\varepsilon^3 \langle f_n'''(z^-), u^{\otimes 3}\rangle$$

for some z^- on the line between x and $x - \varepsilon u$. Adding these two equations gives

$$f_n(x + \varepsilon u) + f_n(x - \varepsilon u) = 2f_n(x) + \varepsilon^2 u^\mathsf{T} f_n''(x) u + \tfrac{1}{6}\varepsilon^3 \langle f_n'''(z^+) - f_n'''(z^-), u^{\otimes 3}\rangle.$$

For any tensor $T \in \mathbb{R}^{D^3}$, $|\langle T, u^{\otimes 3}\rangle| \leq \|T\| \|u^{\otimes 3}\| = \|T\|$, by the Cauchy–Schwarz inequality. Therefore,

$$|u^\mathsf{T} f_n''(x) u| \leq (1/\varepsilon^2)|f_n(x + \varepsilon u) + f_n(x - \varepsilon u) - 2f_n(x)| + \tfrac{1}{3}\varepsilon C.$$

Thus, since (f_n) is pointwise bounded, this implies that $\{u^\mathsf{T} f_n''(x) u : n \in \mathbb{N}\}$ is bounded, for any u with $|u| = 1$. Let $u_1, \ldots, u_k \in \mathbb{R}^D$, with $|u_i| = 1$, such that $u_1 u_1^\mathsf{T}, \ldots, u_k u_k^\mathsf{T}$ is a basis for the vector space $V \subseteq \mathbb{R}^{D \times D}$ of symmetric matrices. (This is possible since $\text{span}\{uu^\mathsf{T} : |u| = 1\} = V$ by the spectral decomposition theorem.) With $\langle A, B\rangle := \sum_{i,j} A_{ij} B_{ij}$, V is an inner product space. Since $\{u_i^\mathsf{T} f_n''(x) u_i : n \in \mathbb{N}\}$ is bounded for each i, and $u_i^\mathsf{T} f_n''(x) u_i = \langle u_i u_i^\mathsf{T}, f_n''(x)\rangle$, then by Lemma S3.3, $\{f_n''(x) : n \in \mathbb{N}\}$ is bounded. Since x is arbitrary, (f_n'') is pointwise bounded. $\qquad \square$

Lemma S3.3. *Suppose V is a finite-dimensional inner product space over $\mathbb{R}$, and let $e_1, \ldots, e_k \in V$ be a basis. If $S \subseteq V$ such that $\{\langle e_i, x \rangle : x \in S\}$ is bounded for each $i = 1, \ldots, k$, then S is bounded.*

Proof. Let G be the Gram matrix of (e_i), i.e., $G_{ij} = \langle e_i, e_j \rangle$. Note that G is positive definite, since for any $a \in \mathbb{R}^k$,

$$a^{\mathsf{T}} G a = \sum_{i,j} a_i a_j G_{ij} = \sum_{i,j} \langle a_i e_i, a_j e_j \rangle = \left\langle \sum_i a_i e_i, \sum_j a_j e_j \right\rangle = \left\| \sum_i a_i e_i \right\|^2 \geq 0, \qquad \text{(S6)}$$

with equality if and only if $\sum_i a_i e_i = 0$, that is, if and only if $a = 0$ (since (e_i) is a linearly independent set). For $x \in V$, define $a(x) \in \mathbb{R}^k$ by the property that $\sum_i a_i(x) e_i = x$ (noting that $a(x)$ always exists and is unique, since (e_i) is a basis). Define $b(x) \in \mathbb{R}^k$ such that $b_i(x) = \langle e_i, x \rangle$. Then for any $x \in V$,

$$b_i(x) = \left\langle e_i, \sum_j a_j(x) e_j \right\rangle = \sum_j a_j(x) \langle e_i, e_j \rangle = \sum_j a_j(x) G_{ij},$$

and thus, $b(x) = G a(x)$. Hence, $a(x) = G^{-1} b(x)$, so by Equation S6,

$$\|x\|^2 = a(x)^{\mathsf{T}} G a(x) = b(x)^{\mathsf{T}} G^{-1} b(x) \leq \|G^{-1}\| |b(x)|^2.$$

By assumption, $\{|b(x)| : x \in S\}$ is bounded, hence, $\{\|x\| : x \in S\}$ is bounded. $\qquad \square$

Lemma S3.4. *Let $E \subseteq \mathbb{R}^D$ be open, convex, and bounded. Let $f_n : E \to \mathbb{R}$ have continuous second derivatives. If $f_n \to f$ pointwise for some $f : E \to \mathbb{R}$, and (f_n'') is uniformly bounded, then f' exists and is continuous, and $f_n' \to f'$ uniformly.*

Proof. First, we show that (f_n') converges pointwise. Let $C = \sup_n \sup_{x \in E} \|f_n''(x)\| < \infty$. Let $x \in E$, and let $\varepsilon > 0$ small enough that $\overline{B_\varepsilon(x)} \subseteq E$. Then for any $u \in \mathbb{R}^D$ with $|u| = 1$, for any m, n, by applying Taylor's theorem to $f_m - f_n$,

$$f_m(x + \varepsilon u) - f_n(x + \varepsilon u) = f_m(x) - f_n(x) + (f_m'(x) - f_n'(x))^{\mathsf{T}} (\varepsilon u) + \tfrac{1}{2}(\varepsilon u)^{\mathsf{T}} (f_m''(z) - f_n''(z))(\varepsilon u)$$

for some z on the line between x and $x + \varepsilon u$. Thus,

$$|(f_m'(x) - f_n'(x))^{\mathsf{T}} u| \leq \frac{1}{\varepsilon} |f_m(x + \varepsilon u) - f_n(x + \varepsilon u)| + \frac{1}{\varepsilon} |f_m(x) - f_n(x)| + \tfrac{1}{2} \varepsilon \|f_m''(z) - f_n''(z)\|.$$

The first two terms on the right go to zero as $m, n \to \infty$ (by pointwise convergence of f_n), and $\|f_m''(z) - f_n''(z)\| \leq \|f_m''(z)\| + \|f_n''(z)\| \leq 2C$, therefore, $\limsup_{m,n \to \infty} |(f_m'(x) - f_n'(x))^{\mathsf{T}} u| \leq \varepsilon C$. Since ε can be arbitrarily small, $|(f_m'(x) - f_n'(x))^{\mathsf{T}} u| \to 0$ as $m, n \to \infty$. Choosing $u = (1, 0, 0, \ldots, 0)$, then $u = (0, 1, 0, \ldots, 0)$, and so on, this implies $|f_m'(x) - f_n'(x)| \to 0$ as $m, n \to \infty$, and hence, $f_n'(x)$ converges.

Next, by Lemma S6.2, (f_n') is equi-Lipschitz, and hence, equicontinuous. Thus, in fact, (f_n') converges uniformly, by Lemma S6.1. Finally, we show that f' exists and $f_n' \to f'$ uniformly; it will follow that f' is continuous, as the limit of a uniformly convergent sequence of continuous functions.

Let $C_{mn} = \sup_{x \in E} |f'_m(x) - f'_n(x)|$. Then $C_{mn} \to 0$ as $m, n \to \infty$, by uniform convergence. To establish the result, it suffices to show that for any $x_0 \in E$, $f'(x_0)$ exists and $f'_n(x_0) \to f'(x_0)$. Fix $x_0 \in E$, and let $B = B_\varepsilon(x_0) \setminus \{x_0\}$ where $\varepsilon > 0$ is small enough that $B \subseteq E$. For $x \in B$, define $\varphi_n(x) = (f_n(x) - f_n(x_0))/|x - x_0|$ and $\varphi(x) = (f(x) - f(x_0))/|x - x_0|$, noting that $\varphi_n \to \varphi$ pointwise. For any $x \in B$, by Taylor's theorem applied to $f_m - f_n$,

$$f_m(x) - f_n(x) = f_m(x_0) - f_n(x_0) + (f'_m(z) - f'_n(z))^\mathsf{T}(x - x_0)$$

for some z on the line between x and x_0, and hence,

$$|\varphi_m(x) - \varphi_n(x)| \leq |f'_m(z) - f'_n(z)| \leq C_{mn} \longrightarrow 0$$

as $m, n \to \infty$. Therefore, $\varphi_n \to \varphi$ uniformly (on B) (by e.g., Rudin, 1976, 7.8).

Now, define $\psi_n(x) = f'_n(x_0)^\mathsf{T}(x - x_0)/|x - x_0|$ and $\psi(x) = v^\mathsf{T}(x - x_0)/|x - x_0|$ for $x \in B$, where $v = \lim_n f'_n(x_0)$. Since $|\psi_n(x) - \psi(x)| \leq |f'_n(x_0) - v| \to 0$ as $n \to \infty$, then $\psi_n \to \psi$ uniformly as well. Hence, $|\varphi_n - \psi_n| \to |\varphi - \psi|$ uniformly (on B).

By the definition of the derivative $f'_n(x_0)$,

$$|\varphi_n(x) - \psi_n(x)| = \frac{|f_n(x) - f_n(x_0) - f'_n(x_0)^\mathsf{T}(x - x_0)|}{|x - x_0|} \xrightarrow[x \to x_0]{} 0.$$

Therefore (by e.g., Rudin, 1976, 7.11),

$$0 = \lim_{n \to \infty} \lim_{x \to x_0} |\varphi_n(x) - \psi_n(x)| = \lim_{x \to x_0} \lim_{n \to \infty} |\varphi_n(x) - \psi_n(x)| = \lim_{x \to x_0} |\varphi(x) - \psi(x)|$$

$$= \lim_{x \to x_0} \frac{|f(x) - f(x_0) - v^\mathsf{T}(x - x_0)|}{|x - x_0|}.$$

Hence, $f'(x_0)$ exists and equals $v = \lim_n f'_n(x_0)$. $\qquad\square$

Proof of Theorem 3.4. First, suppose (f_n) is pointwise bounded. By Lemma S6.2 with $k = 3$, (f''_n) is equi-Lipschitz, and by Lemma S3.2, (f''_n) is pointwise bounded. Thus, since E is bounded, it follows that (f''_n) is uniformly bounded. Therefore, by Lemma S6.2 with $k = 2$, (f'_n) is equi-Lipschitz, and by Lemma S3.1, (f'_n) is pointwise bounded. Thus, likewise, (f'_n) is uniformly bounded. And lastly, applying Lemma S6.2 with $k = 1$, we have that (f_n) is equi-Lipschitz, and hence, uniformly bounded, since it is pointwise bounded by assumption.

Now, assume $f_n \to f$ pointwise. Then in fact, $f_n \to f$ uniformly, by Lemma S6.1, since (f_n) is equi-Lipschitz (as just established), and hence, equicontinuous. By Lemma S3.4, f' exists and $f'_n \to f'$ uniformly. To complete the proof, we show that f'' exists and $f''_n \to f''$ uniformly. For any $i \in \{1, \dots, D\}$, if we define $h_n(x) = f'_n(x)_i$ and $h(x) = f'(x)_i$, then $h_n \to h$ pointwise and (h''_n) is uniformly bounded (since (f'''_n) is uniformly bounded and $\|h''_n(x)\| \leq \|f'''_n(x)\|$); hence, by Lemma S3.4, h' exists and is continuous, and $h'_n \to h'$ uniformly. Since this holds for each coordinate i, then f'' exists, and $f''_n \to f''$ uniformly. $\quad\square$

S4 Proofs of coverage results

Proof of Theorem 4.1. Letting $X_n = -\sqrt{n}(\theta_n - \theta_0)$ and $X \sim Q$,

$$\mathbb{P}(\theta_0 \in S_n) \overset{\text{(a)}}{=} \mathbb{P}(\sqrt{n}(\theta_0 - \theta_n) \in R_n) = \mathbb{P}(X_n \in R_n) \overset{\text{(b)}}{\longrightarrow} \mathbb{P}(X \in R) = Q(R)$$

where step (a) is by the definition of R_n, and (b) is by Lemma 4.2, using assumptions 1 $(X_n \xrightarrow{D} X)$, 3, and 4. To see that $Q(R) = \alpha$, note that $\Pi_n(S_n) \xrightarrow{\text{a.s.}} \alpha$ by assumption and also $\Pi_n(S_n) = Q_n(R_n) \xrightarrow{\text{a.s.}} Q(R)$ since

$$
\begin{aligned}
|Q_n(R_n) - Q(R)| &\leq |Q_n(R_n) - Q(R_n)| + |Q(R_n) - Q(R)| \\
&\leq \sup_{A \in \mathcal{B}} |Q_n(A) - Q(A)| + |Q(R_n) - Q(R)| \xrightarrow{\text{a.s.}} 0
\end{aligned}
$$

by assumption 2 and assumption 3 plus the dominated convergence theorem (Folland, 2013, Theorem 2.24). $\qquad \square$

Proof of Lemma 4.2. For each $k = 1, 2, \ldots$, define $A_k = \{x \in \mathbb{R}^D : d(x, R^c) > 1/k\}$ and $B_k = \{x \in \mathbb{R}^D : d(x, R) \leq 1/k\}$. Note that A_k is open and B_k is closed since $x \mapsto d(x, R)$ and $x \mapsto d(x, R^c)$ are continuous. For any k, by Lemma 4.3 we have that with probability 1, for all n sufficiently large, $A_k \subseteq R_n \subseteq B_k$. Thus, with probability 1, $\liminf_n \big(\mathbb{1}(X_n \in R_n) - \mathbb{1}(X_n \in A_k)\big) \geq \liminf_n \inf_x \big(\mathbb{1}(x \in R_n) - \mathbb{1}(x \in A_k)\big) \geq 0$. It follows that

$$
\liminf_n \mathrm{E}\big(\mathbb{1}(X_n \in R_n) - \mathbb{1}(X_n \in A_k)\big) \geq \mathrm{E} \liminf_n \big(\mathbb{1}(X_n \in R_n) - \mathbb{1}(X_n \in A_k)\big) \geq 0
$$

by Fatou's lemma applied to $\mathbb{1}(X_n \in R_n) - \mathbb{1}(X_n \in A_k) + 1$. (The $+1$ is added to make the function nonnegative, so that Fatou's lemma applies directly.) Therefore, $\liminf_{n \to \infty} \mathbb{P}(X_n \in A_k) \leq \liminf_{n \to \infty} \mathbb{P}(X_n \in R_n)$. Similarly, by reverse Fatou's lemma,

$$
\limsup_n \mathrm{E}\big(\mathbb{1}(X_n \in R_n) - \mathbb{1}(X_n \in B_k)\big) \leq \mathrm{E} \limsup_n \big(\mathbb{1}(X_n \in R_n) - \mathbb{1}(X_n \in B_k)\big) \leq 0,
$$

and therefore, $\limsup_n \mathbb{P}(X_n \in R_n) \leq \limsup_n \mathbb{P}(X_n \in B_k)$. Hence, by the portmanteau theorem (Dudley, 2002, Theorem 11.1.1), for all k,

$$
\begin{aligned}
\mathbb{P}(X \in A_k) &\leq \liminf_n \mathbb{P}(X_n \in A_k) \leq \liminf_n \mathbb{P}(X_n \in R_n) \\
&\leq \limsup_n \mathbb{P}(X_n \in R_n) \leq \limsup_n \mathbb{P}(X_n \in B_k) \leq \mathbb{P}(X \in B_k).
\end{aligned}
$$

Taking limits as $k \to \infty$ and using the fact that $\bigcup_{k=1}^{\infty} A_k = R^\circ$ and $\bigcap_{k=1}^{\infty} B_k = \bar{R}$, we have $\mathbb{P}(X \in R^\circ) = \lim_k \mathbb{P}(X \in A_k) \leq \liminf_n \mathbb{P}(X_n \in R_n) \leq \limsup_n \mathbb{P}(X_n \in R_n) \leq \lim_k \mathbb{P}(X \in B_k) = \mathbb{P}(X \in \bar{R})$ by (Folland, 2013, Theorem 1.8). Further, $\mathbb{P}(X \in R^\circ) = \mathbb{P}(X \in R) = \mathbb{P}(X \in \bar{R})$ since $\mathbb{P}(X \in \partial R) = 0$. Therefore, $\lim_n \mathbb{P}(X_n \in R_n) = \mathbb{P}(X \in R)$. $\qquad \square$

Proof of Lemma 4.3. First, we establish some initial facts. It is straightforward to check that R is convex. R° is nonempty since $m(\bar{R}) \geq m(R) > 0$ and $m(\partial R) = 0$ (Lang, 1986). It follows that R, A, and B are bounded. For any open cube E such that $\bar{E} \subseteq R$, we have $E \subseteq R_n$ for all n sufficiently large, since $\mathbb{1}(x \in R_n) \to \mathbb{1}(x \in R)$ for each corner x of the cube E.

Next, we show that $A \subseteq R_n$ for all n sufficiently large. For each $x \in \bar{A}$, let E_x be a nonempty open cube centered at x such that $\bar{E}_x \subseteq R$. Then $\{E_x : x \in \bar{A}\}$ is an open cover of $\bar{A}$. Since $\bar{A}$ is compact, there is a finite subcover $E_{x_1}, \ldots, E_{x_k}$. Thus, for all n sufficiently large, $A \subseteq \bar{A} \subseteq \bigcup_{i=1}^{k} E_{x_i} \subseteq R_n$.

Now, we show that $R_n \subseteq B$ for all n sufficiently large. Let $S_\delta = \{x \in \mathbb{R}^D : d(x, R) = \delta\}$ for $\delta > 0$. Let $E \subseteq R$ be a nonempty open cube such that $E \subseteq R_n$ for all n sufficiently large. For each $x \in S_{\varepsilon/2}$, define $C_x = \bigcup_{t>1}\{tx + (1-t)z : z \in E\}$. Then C_x is open, as a union of open sets. Note that $y \in C_x$ if and only if $x = sy + (1-s)z$ for some $s \in (0, 1)$, $z \in E$, i.e., if and only if x is a (strict) convex combination of y and some point of E. Thus, $\{C_x : x \in S_{\varepsilon/2}\}$ is an open cover of S_ε (since for any $y \in S_\varepsilon$, the line between y and any $z \in E$ must pass through $S_{\varepsilon/2}$ by the intermediate value theorem applied to $s \mapsto d(sx+(1-s)z, R)$). Since S_ε is compact, there is a finite subcover $C_{x_1}, \ldots, C_{x_k}$ for some $x_1, \ldots, x_k \in S_{\varepsilon/2}$. Since $x_i \in R^c$ for each $i = 1, \ldots, k$, there exists N such that for all $n \geq N$, $x_1, \ldots, x_k \in R_n^c$ and $E \subseteq R_n$. Then for all $n \geq N$, by the convexity of R_n, we have $S_\varepsilon \subseteq \bigcup_{i=1}^k C_{x_i} \subseteq R_n^c$ and hence $R_n \subseteq B$. $\qquad\square$

S5 Proofs for Cox proportional hazards model

Proof of Theorem 7.6. For $\theta \in \mathbb{R}^D$,

$$f_n(\theta) = -\frac{1}{n}\log\mathcal{L}_n^{\mathrm{Cox}}(\theta) - \frac{1}{n}\sum_{i=1}^n Z_i \log n = \frac{1}{n}\sum_{i=1}^n H_{Y_i}^n(\theta)Z_i - \theta^{\mathrm{T}}\left(\tfrac{1}{n}\sum_{i=1}^n X_i Z_i\right)$$

where $H_y^n(\theta) = \log\left(\frac{1}{n}\sum_{j=1}^n \exp(\theta^{\mathrm{T}}X_j)\mathbb{1}(Y_j \geq y)\right)$. Note that f_n is C^∞, as a composition of C^∞ functions. Further, f_n is convex on $\mathbb{R}^D$, since $H_{Y_i}^n(\theta)$ is convex by Lemma S6.6 with $\mu = \frac{1}{n}\sum_{j:Y_j \geq Y_i} \delta_{X_j}$. By Lemma S5.2, $f''(\theta_0)$ is positive definite.

By the strong law of large numbers, $\frac{1}{n}\sum_{i=1}^n X_i Z_i \xrightarrow{\text{a.s.}} \mathrm{E}(XZ)$ as $n \to \infty$, and by Lemma S5.1, for all $\theta \in \mathbb{R}^D$, $\mathrm{E}|h_Y(\theta)Z| < \infty$ and $\frac{1}{n}\sum_{i=1}^n H_{Y_i}^n(\theta)Z_i \xrightarrow{\text{a.s.}} \mathrm{E}(h_Y(\theta)Z)$ as $n \to \infty$. Therefore, for all $\theta \in \mathbb{R}^D$, with probability 1, $f_n(\theta) \to f(\theta)$. Due to convexity, this implies that with probability 1, for all $\theta \in \mathbb{R}^D$, $f_n(\theta) \to f(\theta)$.

Let $m = \sup\{|x| : x \in \mathcal{X}\} < \infty$. Then by Lemma S6.6, $\left|(\partial^3/\partial\theta_j\partial\theta_k\partial\theta_\ell)H_{Y_i}^n(\theta)\right| \leq (2m)^3 = 8m^3$ for all $\theta \in \mathbb{R}^D$. Thus, $\|f_n'''(\theta)\|^2 = \sum_{j,k,\ell}\left|(\partial^3/\partial\theta_j\partial\theta_k\partial\theta_\ell)f_n(\theta)\right|^2 \leq D^3(8m^3)^2$ for all $\theta \in \mathbb{R}^D$, $n \in \mathbb{N}$. Hence, (f_n''') is a.s. uniformly bounded on all of $\mathbb{R}^D$. Thus, for any open ball E containing θ_0, the conditions of Theorem 3.2 are satisfied with probability 1. $\quad\square$

Note that $H_{Y_1}^n(\theta), H_{Y_2}^n(\theta), \ldots$ are not i.i.d., which is why the following lemma is not trivial.

Lemma S5.1. *Suppose $(X, Y, Z), (X_1, Y_1, Z_1), (X_2, Y_2, Z_2), \ldots$ are i.i.d., where $X \in \mathcal{X} \subseteq \mathbb{R}^D$, $Y \geq 0$, and $Z \in \{0, 1\}$. Define $h_y(\theta) = \log\mathrm{E}\left(\exp(\theta^{\mathrm{T}}X)\mathbb{1}(Y \geq y)\right)$ and $H_y^n(\theta) = \log\left(\frac{1}{n}\sum_{j=1}^n \exp(\theta^{\mathrm{T}}X_j)\mathbb{1}(Y_j \geq y)\right)$ for $\theta \in \mathbb{R}^D$, $y \geq 0$. If $\mathcal{X}$ is bounded and the c.d.f. of Y is continuous on $\mathbb{R}$, then for all $\theta \in \mathbb{R}^D$, $\mathrm{E}|h_Y(\theta)Z| < \infty$ and*

$$\frac{1}{n}\sum_{i=1}^n H_{Y_i}^n(\theta)Z_i \xrightarrow[n\to\infty]{\text{a.s.}} \mathrm{E}(h_Y(\theta)Z).$$

Proof. Let $F(y) = \mathbb{P}(Y \leq y)$, $c^* = \sup\{y \in \mathbb{R} : F(y) < 1\}$, and $m = \sup\{|x| : x \in \mathcal{X}\} < \infty$. Since $|X| \leq m$ and F is continuous, $\mathrm{E}|h_Y(\theta)Z| \leq m|\theta| - \mathrm{E}\log(1 - F(Y)) = m|\theta| + 1$ because $F(Y) \sim \mathrm{Uniform}(0, 1)$. Fix $\theta \in \mathbb{R}^D$ and define $g(y) = h_y(\theta)$ and $G_n(y) = H_y^n(\theta)$.

First, we show that for all $c \in (0, c^*)$,

$$\sup_{y \in [0,c]} |G_n(y) - g(y)| \xrightarrow[n \to \infty]{\text{a.s.}} 0. \tag{S7}$$

Let S be a countable dense subset of $[0, c]$ such that $0, c \in S$. For all $y \in S$, $G_n(y) \xrightarrow{\text{a.s.}} g(y) \in \mathbb{R}$ by the strong law of large numbers since $0 < \mathrm{E}(e^{\theta^{\mathrm{T}} X} \mathbb{1}(Y \geq y)) \leq e^{m|\theta|} < \infty$. Next, G_n is a non-increasing function on $[0, c]$ (that is, if $0 \leq y < y' \leq c$ then $G_n(y) \geq G_n(y')$) since $y \mapsto \mathbb{1}(Y_j \geq y)$ is non-increasing. Further, $g(y)$ is continuous on $[0, c]$ by the dominated convergence theorem, since $|e^{\theta^{\mathrm{T}} X} \mathbb{1}(Y \geq y)| \leq e^{m|\theta|}$ and $\mathbb{P}(Y = y) = 0$ by the continuity of F. Thus, with probability 1, for all n sufficiently large, G_n is finite on $[0, c]$ since $G_n(0) \xrightarrow{\text{a.s.}} g(0)$ and $G_n(c) \xrightarrow{\text{a.s.}} g(c)$. It follows that $\sup_{y \in [0,c]} |G_n(y) - g(y)| \xrightarrow{\text{a.s.}} 0$ by Lemma S6.4.

Second, we show that for all $c \in (0, c^*)$,

$$\frac{1}{n} \sum_{i=1}^{n} G_n(Y_i) Z_i \mathbb{1}(Y_i \leq c) \xrightarrow[n \to \infty]{\text{a.s.}} \mathrm{E}\big(g(Y) Z \mathbb{1}(Y \leq c)\big). \tag{S8}$$

To see this, observe that by Equation S7,

$$\left| \frac{1}{n} \sum_{i=1}^{n} G_n(Y_i) Z_i \mathbb{1}(Y_i \leq c) - \frac{1}{n} \sum_{i=1}^{n} g(Y_i) Z_i \mathbb{1}(Y_i \leq c) \right|$$

$$\leq \frac{1}{n} \sum_{i=1}^{n} |G_n(Y_i) - g(Y_i)| \mathbb{1}(Y_i \leq c) \leq \sup_{y \in [0,c]} |G_n(y) - g(y)| \xrightarrow[n \to \infty]{\text{a.s.}} 0$$

and $\frac{1}{n} \sum_{i=1}^{n} g(Y_i) Z_i \mathbb{1}(Y_i \leq c) \xrightarrow[n \to \infty]{\text{a.s.}} \mathrm{E}\big(g(Y) Z \mathbb{1}(Y \leq c)\big)$ by the strong law of large numbers.

Third, we show that for all $c \in (0, c^*)$,

$$\limsup_{n \to \infty} \left| \frac{1}{n} \sum_{i=1}^{n} G_n(Y_i) Z_i \mathbb{1}(Y_i > c) \right| \overset{\text{a.s.}}{\leq} m|\theta| p_c - p_c \log p_c + p_c \tag{S9}$$

where $p_c = \mathbb{P}(Y > c)$. This follows from the fact that

$$\left| \frac{1}{n} \sum_{i=1}^{n} G_n(Y_i) Z_i \mathbb{1}(Y_i > c) \right| \leq \frac{1}{n} \sum_{i=1}^{n} |G_n(Y_i)| \mathbb{1}(Y_i > c)$$

$$\leq \frac{1}{n} \sum_{i=1}^{n} \left(m|\theta| - \log \left(\frac{1}{n} \sum_{j=1}^{n} \mathbb{1}(Y_j \geq Y_i) \right) \right) \mathbb{1}(Y_i > c)$$

$$\overset{\text{a.s.}}{=} m|\theta| K_n / n - \frac{1}{n} \sum_{k=1}^{K_n} \log(k/n)$$

$$\xrightarrow[n \to \infty]{\text{a.s.}} m|\theta| p_c - \int_0^{p_c} (\log x) dx = m|\theta| p_c - p_c \log p_c + p_c$$

where $K_n = \sum_{i=1}^{n} \mathbb{1}(Y_i > c)$, using that $\mathbb{P}(Y_i = Y_j) = 0$ for $i \neq j$ by continuity of F.

Now, we put these pieces together to obtain the result. Writing $\frac{1}{n}\sum_{i=1}^{n}G_n(Y_i)Z_i = \frac{1}{n}\sum_{i=1}^{n}G_n(Y_i)Z_i\mathbb{1}(Y_i \leq c) + \frac{1}{n}\sum_{i=1}^{n}G_n(Y_i)Z_i\mathbb{1}(Y_i > c)$, for all $c \in (0, c^*)$ we have

$$\left|\frac{1}{n}\sum_{i=1}^{n}G_n(Y_i)Z_i - \mathrm{E}(g(Y)Z)\right| \leq \left|\frac{1}{n}\sum_{i=1}^{n}G_n(Y_i)Z_i\mathbb{1}(Y_i \leq c) - \mathrm{E}\big(g(Y)Z\mathbb{1}(Y \leq c)\big)\right|$$
$$+ \left|\mathrm{E}\big(g(Y)Z\mathbb{1}(Y \leq c)\big) - \mathrm{E}\big(g(Y)Z\big)\right| + \left|\frac{1}{n}\sum_{i=1}^{n}G_n(Y_i)Z_i\mathbb{1}(Y_i > c)\right|,$$

and therefore, by Equations S8 and S9,

$$\limsup_{n\to\infty}\left|\frac{1}{n}\sum_{i=1}^{n}G_n(Y_i)Z_i - \mathrm{E}(g(Y)Z)\right|$$
$$\overset{\text{a.s.}}{\leq} \left|\mathrm{E}\big(g(Y)Z\mathbb{1}(Y \leq c)\big) - \mathrm{E}(g(Y)Z)\right| + m|\theta|p_c - p_c\log p_c + p_c. \tag{S10}$$

Let $c_1, c_2, \ldots \in (0, c^*)$ such that $c_k \to c^*$. Then $p_{c_k} \to p_{c^*} = 0$ by continuity of F, and thus, $m|\theta|p_{c_k} - p_{c_k}\log p_{c_k} + p_{c_k} \to 0$ as $k \to \infty$. Further, $\mathrm{E}\big(g(Y)Z\mathbb{1}(Y \leq c_k)\big) \to \mathrm{E}(g(Y)Z)$ by the dominated convergence theorem, since $|g(Y)Z\mathbb{1}(Y \leq c_k)| \leq |g(Y)Z|$, $\mathrm{E}|g(Y)Z| < \infty$, and $\mathbb{1}(Y \leq c_k) \overset{\text{a.s.}}{\to} 1$ as $k \to \infty$. Applying Equation S10 to each c_k and taking limits as $k \to \infty$, we have that $\limsup_{n\to\infty}\left|\frac{1}{n}\sum_{i=1}^{n}G_n(Y_i)Z_i - \mathrm{E}(g(Y)Z)\right| = 0$ almost surely. $\square$

Lemma S5.2. *Under the conditions of Theorem 7.6, $f''(\theta)$ is positive definite for all $\theta \in \mathbb{R}^D$.*

Proof. Recall that $f(\theta) = \mathrm{E}\big(h_Y(\theta)Z\big) - \theta^\mathsf{T}\mathrm{E}(XZ)$ where $h_y(\theta) = \log\mathrm{E}(e^{\theta^\mathsf{T}X}\mathbb{1}(Y \geq y))$ for $\theta \in \mathbb{R}^D$. First, we put $h_y(\theta)$ in the form of $\kappa(\theta)$ in Lemma S6.6 by noting that $h_y(\theta) = \log\mathrm{E}(e^{\theta^\mathsf{T}X}\mathbb{P}(Y \geq y \mid X)) = \log\int\exp(\theta^\mathsf{T}x)\mu_y(dx)$ where $\mu_y(dx) = \mathbb{P}(Y \geq y \mid X = x)P(dx)$ and P is the distribution of X (Dudley, 2002, 10.2.1-10.2.2). Let $m = \sup\{|x| : x \in \mathcal{X}\} < \infty$. We have $|h_y(\theta)| < \infty$ for all $\theta \in \mathbb{R}^D$ and all $y \geq 0$ because $\exp(-m|\theta|) \leq \exp(\theta^\mathsf{T}X) \leq \exp(m|\theta|)$, and thus $-\infty < -m|\theta| + \log\mathbb{P}(Y \geq y) \leq h_y(\theta) \leq m|\theta| + \log\mathbb{P}(Y \geq y) < \infty$ due to assumptions 1 and 4 of Theorem 7.6.

For any given $\theta \in \mathbb{R}^D$ and $y \geq 0$, following Lemma S6.6, we define a probability measure $\tilde{P} = \tilde{P}_{\theta,y}$ on $\mathcal{X}$ by $\tilde{P}(dx) = \exp(\theta^\mathsf{T}x - h_y(\theta))\mathbb{P}(Y \geq y \mid X = x)P(dx)$. Note that P and $\tilde{P}$ are mutually absolutely continuous since $\exp(\theta^\mathsf{T}x - h_y(\theta))\mathbb{P}(Y \geq y \mid X = x)$ is strictly positive for all $x \in \mathcal{X}$. By Lemma S6.6, $h_y'(\theta) = \mathrm{E}(\tilde{X})$ and $h_y''(\theta) = \mathrm{Cov}(\tilde{X})$ where $\tilde{X} \sim \tilde{P}$. We claim that for any nonzero $a \in \mathbb{R}^D$, $a^\mathsf{T}h_y''(\theta)a > 0$. To see this, suppose $a \in \mathbb{R}^D$ such that $a^\mathsf{T}h_y''(\theta)a = 0$. Since $a^\mathsf{T}h_y''(\theta)a = \mathrm{Var}(a^\mathsf{T}\tilde{X})$, it follows that $\mathbb{P}(a^\mathsf{T}\tilde{X} = \mathrm{E}(a^\mathsf{T}\tilde{X})) = 1$. But then $\mathbb{P}(a^\mathsf{T}X = \mathrm{E}(a^\mathsf{T}\tilde{X})) = 1$ since $P \ll \tilde{P}$. Hence, $a^\mathsf{T}X$ is a.s. equal to a constant, so $\mathrm{Var}(a^\mathsf{T}X) = 0$, which implies $a = 0$ by assumption 3 of Theorem 7.6.

To justify differentiating under the expectation in $\mathrm{E}(h_Y(\theta)Z)$, we apply Folland (2013, Theorem 2.27b) using the following bounds. First, $\mathrm{E}|h_Y(\theta)Z| < \infty$ by Lemma S5.1. Next, $|\tilde{X}| \leq m$ because $\tilde{P}$ is supported on $\mathcal{X}$. Thus, $|\frac{\partial}{\partial\theta_j}h_y(\theta)z| = |\mathrm{E}(\tilde{X}_j)z| \leq \mathrm{E}|\tilde{X}_j| \leq \mathrm{E}|\tilde{X}| \leq m$ and $|\frac{\partial^2}{\partial\theta_j\partial\theta_k}h_y(\theta)z| = |\mathrm{Cov}(\tilde{X}_j, \tilde{X}_k)z| \leq \mathrm{E}|\tilde{X}_j||\tilde{X}_k| + \mathrm{E}|\tilde{X}_j|\mathrm{E}|\tilde{X}_k| \leq 2m^2$ for $z \in \{0, 1\}$.

Hence, $f''(\theta) = \mathrm{E}\big(h_Y''(\theta)Z\big)$, and we have that for any nonzero $a \in \mathbb{R}^D$, $a^\mathsf{T}f''(\theta)a = \mathrm{E}\big(a^\mathsf{T}h_Y''(\theta)aZ\big) > 0$ because $a^\mathsf{T}h_Y''(\theta)a > 0$ and $\mathbb{P}(Z = 1) > 0$ due to assumption 3 of Theorem 7.6. Therefore, $f''(\theta)$ is positive definite. $\square$

S6 Supporting results

This section contains miscellaneous supporting results used in the proofs. A metric space E is *totally bounded* if for any $\delta > 0$, there exist $x_1, \ldots, x_k \in E$, for some $k \in \mathbb{N}$, such that $E = \bigcup_{i=1}^{k} \{x \in E : d(x, x_i) < \delta\}$. In particular, any bounded subset of a Euclidean space is totally bounded.

Lemma S6.1. *Suppose $h_n : E \to F$ for $n \in \mathbb{N}$, where E is a totally bounded metric space and F is a normed space. If (h_n) converges pointwise and is equicontinuous, then it converges uniformly.*

Proof. Let $\varepsilon > 0$. Choose $\delta > 0$ by equicontinuity, so that for any $n \in \mathbb{N}$, $x, y \in E$, if $d(x, y) < \delta$ then $\|h_n(x) - h_n(y)\| < \varepsilon$. Choose $x_1, \ldots, x_k \in E$ by totally boundedness, and by pointwise convergence, let N such that for all $m, n > N$, for all $i \in \{1, \ldots, k\}$, $\|h_m(x_i) - h_n(x_i)\| < \varepsilon$. Then, for any $x \in E$, there is some $i \in \{1, \ldots, k\}$ such that $d(x, x_i) < \delta$, and thus

$$\|h_m(x) - h_n(x)\| \le \|h_m(x) - h_m(x_i)\| + \|h_m(x_i) - h_n(x_i)\| + \|h_n(x_i) - h_n(x)\| < 3\varepsilon$$

for any $m, n > N$. Therefore, (h_n) converges uniformly (by e.g., Rudin, 1976, 7.8). $\qquad \square$

When all the kth order partial derivatives of f exist, let $f^{(k)}(x)$ denote the k-way tensor of kth derivatives; in particular, $f^{(1)} = f'$, $f^{(2)} = f''$, and so on. When these derivatives are continuous, the order of differentiation does not matter (Rudin, 1976, exercise 9.29).

Lemma S6.2. *Let $E \subseteq \mathbb{R}^D$ be open and convex, and let $f_n : E \to \mathbb{R}$ for $n \in \mathbb{N}$. For any $k \in \mathbb{N}$, if each f_n has continuous kth-order derivatives and $(f_n^{(k)})$ is uniformly bounded, then $(f_n^{(k-1)})$ is equi-Lipschitz.*

Proof. First, we prove the case of $k = 1$. Let $C = \sup_n \sup_{x \in E} |f_n'(x)| < \infty$. By Taylor's theorem, for any $n \in \mathbb{N}$, $x, y \in E$, $f_n(x) = f_n(y) + f_n'(z)^\mathsf{T}(x - y)$ for some z on the line between x and y, and therefore,

$$|f_n(x) - f_n(y)| \le |f_n'(z)| \, |x - y| \le C|x - y|.$$

Thus, (f_n) is equi-Lipschitz.

For notational clarity, we prove the case of $k = 3$, and observe that the extension from this to the general case is immediate. For any $i, j \in \{1, \ldots, D\}$, if we define $h_n(x) = f_n''(x)_{ij} = \frac{\partial^2}{\partial x_i \partial x_j} f_n(x)$, then (h_n') is uniformly bounded (since $|h_n'(x)| \le \|f_n'''(x)\|$ and (f_n''') is uniformly bounded), and hence, (h_n) is equi-Lipschitz by the case of $k = 1$ just proven. Thus, (f_n'') is equi-Lipschitz, since if C_{ij} is the equi-Lipschitz constant for entry (i, j), then

$$\|f_n''(x) - f_n''(y)\|^2 = \sum_{i,j} |f_n''(x)_{ij} - f_n''(y)_{ij}|^2 \le C^2 |x - y|^2$$

where $C^2 = \sum_{i,j} C_{ij}^2$. $\qquad \square$

Lemma S6.3. *Let $B \subseteq \mathbb{R}^D$ be open and let $f : B \to \mathbb{R}$ be differentiable. If $x_0 \in B$ such that $f(x) \ge f(x_0)$ for all $x \in B$, then $f'(x_0) = 0$.*

Proof. For any $u \in \mathbb{R}^D$ with $|u| = 1$, $f'(x_0)^\mathsf{T} u = \lim_{\varepsilon \to 0}(f(x_0 + \varepsilon u) - f(x_0)) \geq 0$. If $f'(x_0) \neq 0$, then choosing $u = -f'(x_0)/|f'(x_0)|$, we have $0 \leq f'(x_0)^\mathsf{T} u = -|f'(x_0)| < 0$, a contradiction. $\qquad\square$

Lemma S6.4. *Let $a, b \in \mathbb{R}$ such that $a < b$, let $g : [a, b] \to \mathbb{R}$ be continuous, and for $n \in \mathbb{N}$, let $g_n : [a, b] \to \mathbb{R}$ be a non-increasing function. If there is a dense subset $S \subseteq [a, b]$ such that $a, b \in S$ and $g_n(y) \to g(y)$ for all $y \in S$, then $\sup_{y \in [a,b]} |g_n(y) - g(y)| \longrightarrow 0$ as $n \to \infty$.*

Lemma S6.4 is straightforward to verify, so we omit the proof. Lemmas S6.5 and S6.6 are standard well-known results, but we provide precise statements and proofs for completeness. We write S° to denote the interior of S.

Lemma S6.5. *Let μ be a Borel measure on $\mathbb{R}^D$ and define $G(\theta) = \int_{\mathbb{R}^D} \exp(\theta^\mathsf{T} x)\mu(dx)$ for $\theta \in \mathbb{R}^D$. Let $S = \{\theta \in \mathbb{R}^D : G(\theta) < \infty\}$. Then G is C^∞ on S° and for all $\theta \in S^\circ$, $k \in \{0, 1, 2, \ldots\}$, $i_1, \ldots, i_k \in \{1, \ldots, D\}$, we have*

$$\frac{\partial}{\partial \theta_{i_1}} \cdots \frac{\partial}{\partial \theta_{i_k}} G(\theta) = \int x_{i_1} \cdots x_{i_k} \exp(\theta^\mathsf{T} x)\mu(dx). \tag{S11}$$

Proof. We proceed by induction. By construction, for all $\theta \in S^\circ$, $\int |e^{\theta^\mathsf{T} x}|\mu(dx) < \infty$ and Equation S11 holds when $k = 0$. Fix $i_1, \ldots, i_k \in \{1, \ldots, D\}$ and suppose that for all $\theta \in S^\circ$, $\int |x_{i_1} \cdots x_{i_k} e^{\theta^\mathsf{T} x}|\mu(dx) < \infty$ and Equation S11 holds. Let $j \in \{1, \ldots, D\}$ and $\theta_0 \in S^\circ$. Define $u = (0, \ldots, 0, 1, 0, \ldots, 0) \in \mathbb{R}^D$ where the 1 is in the jth position. Choose $\varepsilon > 0$ such that $\theta_0 + tu \in S^\circ$ for all $t \in [-2\varepsilon, 2\varepsilon]$. Define $f(x, t) = x_{i_1} \cdots x_{i_k} e^{(\theta_0 + tu)^\mathsf{T} x}$ and $F(t) = \int f(x, t)\mu(dx)$ for $x \in \mathbb{R}^D$, $t \in [-2\varepsilon, 2\varepsilon]$. Note that $\int |f(x, t)|\mu(dx) < \infty$ for all $t \in [-2\varepsilon, 2\varepsilon]$ by the induction hypothesis. Define $g(x) = |f(x, 2\varepsilon)|/\varepsilon + |f(x, -2\varepsilon)|/\varepsilon$. It is straightforward to verify that $|\frac{\partial f}{\partial t}(x, t)| = |x_j f(x, t)| \leq g(x)$ for all $x \in \mathbb{R}^D$, $t \in [-\varepsilon, \varepsilon]$, by using the inequality $|x_j| \leq e^{\varepsilon|x_j|}/\varepsilon$. Further, $\int |g(x)|\mu(dx) < \infty$ by the induction hypothesis. Therefore, F is differentiable and $F'(t) = \int \frac{\partial f}{\partial t}(x, t)\mu(dx)$ for all $t \in (-\varepsilon, \varepsilon)$ by Folland (2013, Theorem 2.27b).

Putting these pieces together, we have

$$\frac{\partial}{\partial \theta_j}\bigg|_{\theta = \theta_0} \frac{\partial}{\partial \theta_{i_1}} \cdots \frac{\partial}{\partial \theta_{i_k}} G(\theta) = \frac{\partial}{\partial \theta_j}\bigg|_{\theta = \theta_0} \int x_{i_1} \cdots x_{i_k} \exp(\theta^\mathsf{T} x)\mu(dx)$$

$$= \frac{\partial}{\partial t}\bigg|_{t=0} \int f(x, t)\mu(dx) = F'(0) = \int \frac{\partial f}{\partial t}(x, 0)\mu(dx)$$

$$= \int x_j f(x, 0)\mu(dx) = \int x_j x_{i_1} \cdots x_{i_k} \exp(\theta_0^\mathsf{T} x)\mu(dx)$$

and $\int |x_j x_{i_1} \cdots x_{i_k} e^{\theta_0^\mathsf{T} x}|\mu(dx) = \int |\frac{\partial f}{\partial t}(x, 0)|\mu(dx) \leq \int |g(x)|\mu(dx) < \infty$. Since $j \in \{1, \ldots, D\}$ and $\theta_0 \in S^\circ$ are arbitrary, this completes the induction step. $\qquad\square$

Lemma S6.6. *Let μ be a Borel measure on $\mathbb{R}^D$ and define $\kappa(\theta) = \log \int_{\mathbb{R}^D} \exp(\theta^\mathsf{T} x)\mu(dx)$ for $\theta \in \mathbb{R}^D$. Let $\Theta = \{\theta \in \mathbb{R}^D : |\kappa(\theta)| < \infty\}$, and define $P_\theta(A) = \int_A \exp(\theta^\mathsf{T} x - \kappa(\theta))\mu(dx)$ for $\theta \in \Theta$ and $A \subseteq \mathbb{R}^D$ Borel measurable. Then Θ is a convex set and κ is convex on Θ. Further, for all θ in the interior of Θ, for all $i, j, k \in \{1, \ldots, D\}$, if $X \sim P_\theta$ then*

1. $\dfrac{\partial \kappa}{\partial \theta_i}(\theta) = \mathrm{E}(X_i)$,

2. $\dfrac{\partial^2 \kappa}{\partial \theta_i \partial \theta_j}(\theta) = \mathrm{E}\big((X_i - \mathrm{E}X_i)(X_j - \mathrm{E}X_j)\big) = \mathrm{Cov}(X_i, X_j)$, and

3. $\dfrac{\partial^3 \kappa}{\partial \theta_i \partial \theta_j \partial \theta_k}(\theta) = \mathrm{E}\big((X_i - \mathrm{E}X_i)(X_j - \mathrm{E}X_j)(X_k - \mathrm{E}X_k)\big)$.

More succinctly, items 1 and 2 state that $\kappa'(\theta) = \mathrm{E}(X)$ and $\kappa''(\theta) = \mathrm{Cov}(X)$ where $X \sim P_\theta$.

Proof. Convexity of Θ and κ is a straightforward application of Hölder's inequality. Define $G(\theta) = \int \exp(\theta^{\mathsf{T}} x)\mu(dx)$ for $\theta \in \mathbb{R}^D$. By Lemma S6.5, G is C^∞ on the interior of Θ and its partial derivatives are given by Equation S11. The identities in items 1-3 are straightforward to derive using Equation S11 and the chain rule. $\qquad\square$